

从自动化到自治

-Oracle 优化案例实践和 19c 新特性揭秘

盖国强

云和恩墨 创始人，首席架构师

想做团队的领跑者 需要迈过这些“槛”

成长型企业，易忽视人才体系化培养
企业转型加快，团队能力又跟不上

VS

从基础到进阶，超100+一线实战
技术专家带你系统化学习成长

团队成员技能水平不一，
难以一“敌”百人需求

VS

解决从小白到资深技术人所遇到
80%的问题

寻求外部培训，奈何价更高且
集中式学习

VS

多样、灵活的学习方式，包括
音频、图文 和视频

学习效果难以统计，产生不良循环

VS

获取员工学习报告，查看学习
进度，形成闭环



课程顾问「橘子」

回复「QCon」
免费获取
学习解决方案



极客时间企业账号 # 解决技术人成长路上的学习问题

自我介绍



盖国强 云和恩墨信息技术有限公司 创始人

- 国内第一个Oracle ACE及ACE总监;
- 致力于技术分享与传播
 - 自2001年, 技术论坛ITPUB的主要倡导者之一;
 - 自2003年, 已经编辑、编译、著作出版了14本技术书籍;
 - 自2010年, 和张乐奕共创 Oracle用户组 - ACOUG,开展持续的公益活动;
- 走在技术创业的道路上
 - 自2008年, 尝试单飞与创业;
 - 自2011年, 和创业伙伴发起成立云和恩墨;
 - 自2019年, 公司伙伴超过500人;



目 录

- Oracle 的变革之路
- 性能优化与数据库进化
- Oracle 19c新特性

Oracle 的云上变革之路 – 全面转向



Oracle Cloud – Business Summary 2016

- SaaS and PaaS: Sold more than anybody last fiscal year
- SaaS and PaaS: Will sell more (\$2+ billion) than anybody this year
- SaaS and PaaS: Growing faster than anybody – 82% last quarter
- Cloud Revenue: Almost \$1 Billion last quarter
- IaaS: Gen2 Technology Leapfrogs Amazon – Huge Opportunity

ORACLE

Copyright © 2018, Oracle and/or its affiliates. All rights reserved. |

World's First Cloud Native Management & Security System



- **Complete and Integrated System**
 - Monitor and analyze ALL users and assets in a single system
- **Powered by Machine Learning (ML)**
 - ML-based insights and anomaly detection
- **Automated Remediation**
 - Automated operational workflows and real-time security remediation

A New Generation of Cloud Computing

New Hardware – New Software – Advanced Cyber Security Built-in


- **Impenetrable Barrier:** Network of Dedicated Cloud Control Computers
 - Barrier: Protects Cloud Perimeter and Users Zones
 - Impenetrable: No User Access to Cloud Control Computers and Memory
- **Autonomous RoBots:** Machine Learning RoBots Detect and Kill Threats
 - Database Immediately Patches Itself while Running – **Stops Data Theft**
 - No Delay for Human Process or Downtime Window
- **Autonomous Database:** Nothing to Learn, Nothing to Do
 - Eliminates Human Errors, Tolerates Hardware and Software Errors
 - Lowest Labor Costs, Lowest Operational Costs



Oracle 的云上变革之路 – 数据为王

Google Cloud's new CEO is executing the playbook that Larry Ellison apparently wouldn't let him run with at Oracle

Rosalie Chan Apr. 12, 2019, 2:11 PM



Google is supporting customer decisions to use the cloud provider of their choice, even if it's a rival company.

Google Acquires Aloomo, Cloud Database Migration Software Tools

Google Cloud Platform (GCP) acquires Aloomo to counter AWS Aurora & Redshift; Microsoft Azure SQL Data Warehouse; and Oracle Cloud Autonomous Data Warehouse.

by Joe Panettieri • Feb 19, 2019

Google has acquired Aloomo, a key move that could boost [Google Cloud Platform \(GCP\)](#) as a data warehousing system vs. [Amazon Web Services \(AWS\) Aurora](#) and [AWS Redshift](#); [Microsoft Azure SQL Data Warehouse](#) and [Oracle Cloud Autonomous Data Warehouse](#). Financial terms of the deal were not disclosed.

The move could turn heads in the global systems integrator market, where large partners are helping customers to move open- and closed-source databases to public clouds. Smaller MSPs moving into the managed database services market may also be intrigued.

Aloomo develops migration tools and integrations that allow customers to move their data from multiple sources to a single data warehouse. The company is well-known in the [ETL market](#). ETL is short for extract, transform, load — three database functions that are combined into one tool to pull data out of one database and place it into another database. The company launched in 2013 and had raised \$15 million in funding, according to [Crunchbase](#).




Aloomo Co-founder Yoni Broyde




Aloomo Co-founder Yair Weinberger


云数据库 – Oracle 12c 的核心特性



多租户: Massive Cost Savings
and Cloud Agility with
Multitenant



内存存储: Massive
Performance
with **Database In-Memory**



数据库分片: Massive Web
Applications with
Sharding

数据库整合集中、OLTP 和 OLAP 一体、以及大规模并发的负载分散。

自治数据库 – Oracle 18c 的云上部署

18^c

供应 - Provision

Create RAC cluster with
Data Guard Standby

安全 - Secure

Encrypt data, Database
Vault, apply security
patches online

更新 - Update

Online patching and
upgrade of database

保护 - Protect

Backup, failover, repair

整合 - Consolidation

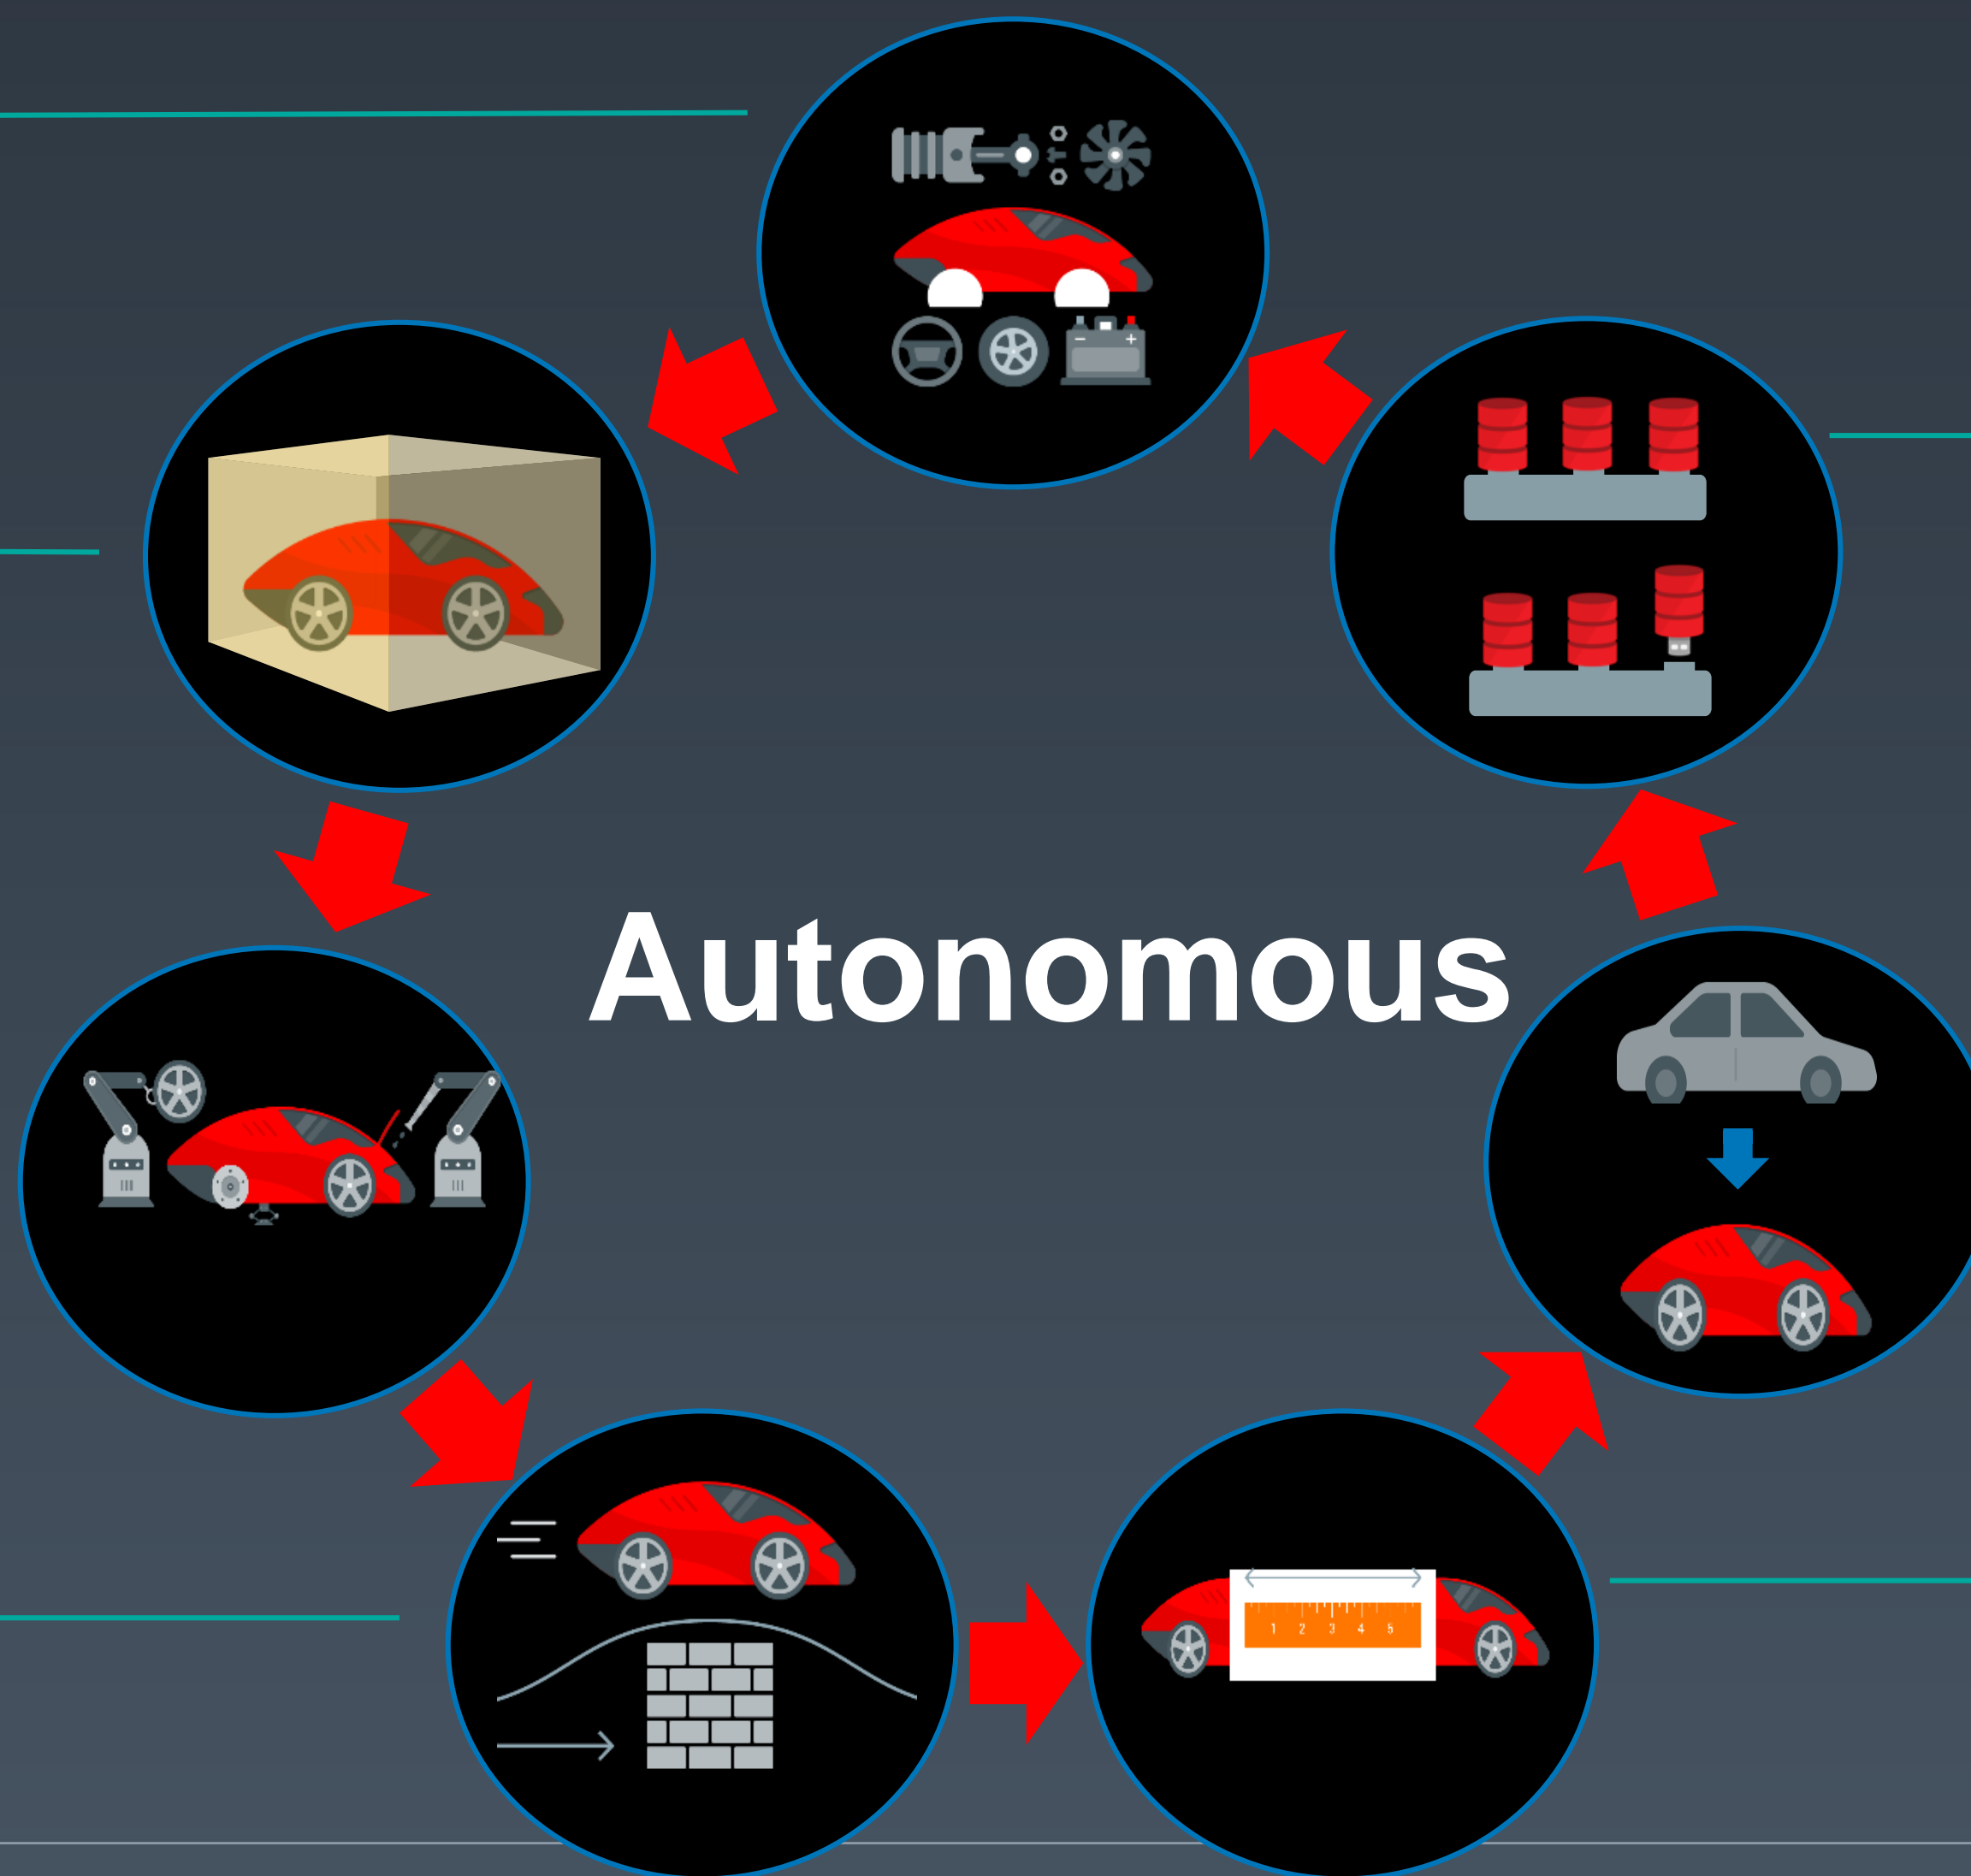
User driven PDB and
CDB creation

迁移 - Migrate

Easy DB migration,
load data from object
store

扩展 - Scale

Elastically adjust
OCPUs, expand DB



自治数据库 – Oracle 19c 的机器学习



索引 - Indexes

Automatic creation and monitoring of indexes

统计信息 - Statistics

Real-time maintenance of optimizer statistics

执行计划 - Plans

Automatic detection and correction of regressions due to plan changes

控制SQL - Runaway SQL
Automatically kill runaway SQL and prevent re-execution until tuned

配置 - Configuration
Automatically tune memory, process, sessions, etc.

Machine Learning

目 录

- Oracle 的变革之路
- **性能优化与数据库进化**
- Oracle 19c新特性

我的成长：曾经年少做开发

```
SELECT "SP_TRANS"."TRANS_NO", ... "SP_ITEM"."CHART_ID", "SP_ITEM"."SPECIFICATION"
FROM  "SP_TRANS", "SP_TRANS_SUB", "SP_CHK", "SP_CHK_SUB",
      "SP_RECEIVE", "SP_RECEIVE_SUB", "SP_ITEM"
WHERE ( "SP_TRANS_SUB"."TRANS_NO" = "SP_TRANS"."TRANS_NO" ) and
      ("SP_TRANS"."BILL_NO" = "SP_CHK"."CHK_NO") and
      ( "SP_CHK_SUB"."CHK_NO" = "SP_CHK"."CHK_NO" ) and
      ( "SP_CHK"."RECEIVE_NO" = "SP_RECEIVE"."RECEIVE_NO" ) and
      ( "SP_RECEIVE_SUB"."RECEIVE_NO" = "SP_RECEIVE"."RECEIVE_NO" ) and
      .....
      ( "SP_CHK_SUB"."COUNTRY" = "SP_RECEIVE_SUB"."COUNTRY" ) and
      ( "SP_CHK_SUB"."PLAN_NO" = "SP_RECEIVE_SUB"."PLAN_NO" ) and
      ( "SP_CHK_SUB"."PLAN_LINE" = "SP_RECEIVE_SUB"."PLAN_LINE" ) and
      (to_char("SP_TRANS"."TRANSDATE" , 'YYYY-MM-DD') >= '2003-01-01');
```

130 rows selected.

Elapsed: 00: 29: 1785.47



SQL背后的世界：误入歧途DBA

```
SELECT /*+ ordered */ "SP_TRANS"."TRANS_NO", ... "SP_ITEM"."CHART_ID","SP_ITEM"."SPECIFICATION"

FROM  "SP_TRANS", "SP_TRANS_SUB", "SP_CHK", "SP_CHK_SUB",

      "SP_RECEIVE", "SP_RECEIVE_SUB", "SP_ITEM"

WHERE ( "SP_TRANS_SUB"."TRANS_NO" = "SP_TRANS"."TRANS_NO" ) and

      ("SP_TRANS"."BILL_NO" = "SP_CHK"."CHK_NO") and

      ( "SP_CHK_SUB"."CHK_NO" = "SP_CHK"."CHK_NO" ) and

      ( "SP_CHK"."RECEIVE_NO" = "SP_RECEIVE"."RECEIVE_NO" ) and

      ( "SP_RECEIVE_SUB"."RECEIVE_NO" = "SP_RECEIVE"."RECEIVE_NO" ) and

      .....

      ( "SP_CHK_SUB"."COUNTRY" = "SP_RECEIVE_SUB"."COUNTRY" ) and

      ( "SP_CHK_SUB"."PLAN_NO" = "SP_RECEIVE_SUB"."PLAN_NO" ) and

      ( "SP_CHK_SUB"."PLAN_LINE" = "SP_RECEIVE_SUB"."PLAN_LINE" ) and

      (to_char("SP_TRANS"."TRANSDATE" , 'YYYY-MM-DD') >= '2003-01-01') ;
```

130 rows selected.

Elapsed: 00: 00: 05.67

SQL背后的世界：多表如何联合

1 SP_TRANS

2 SP_TRANS_SUB

3 SP_CHK

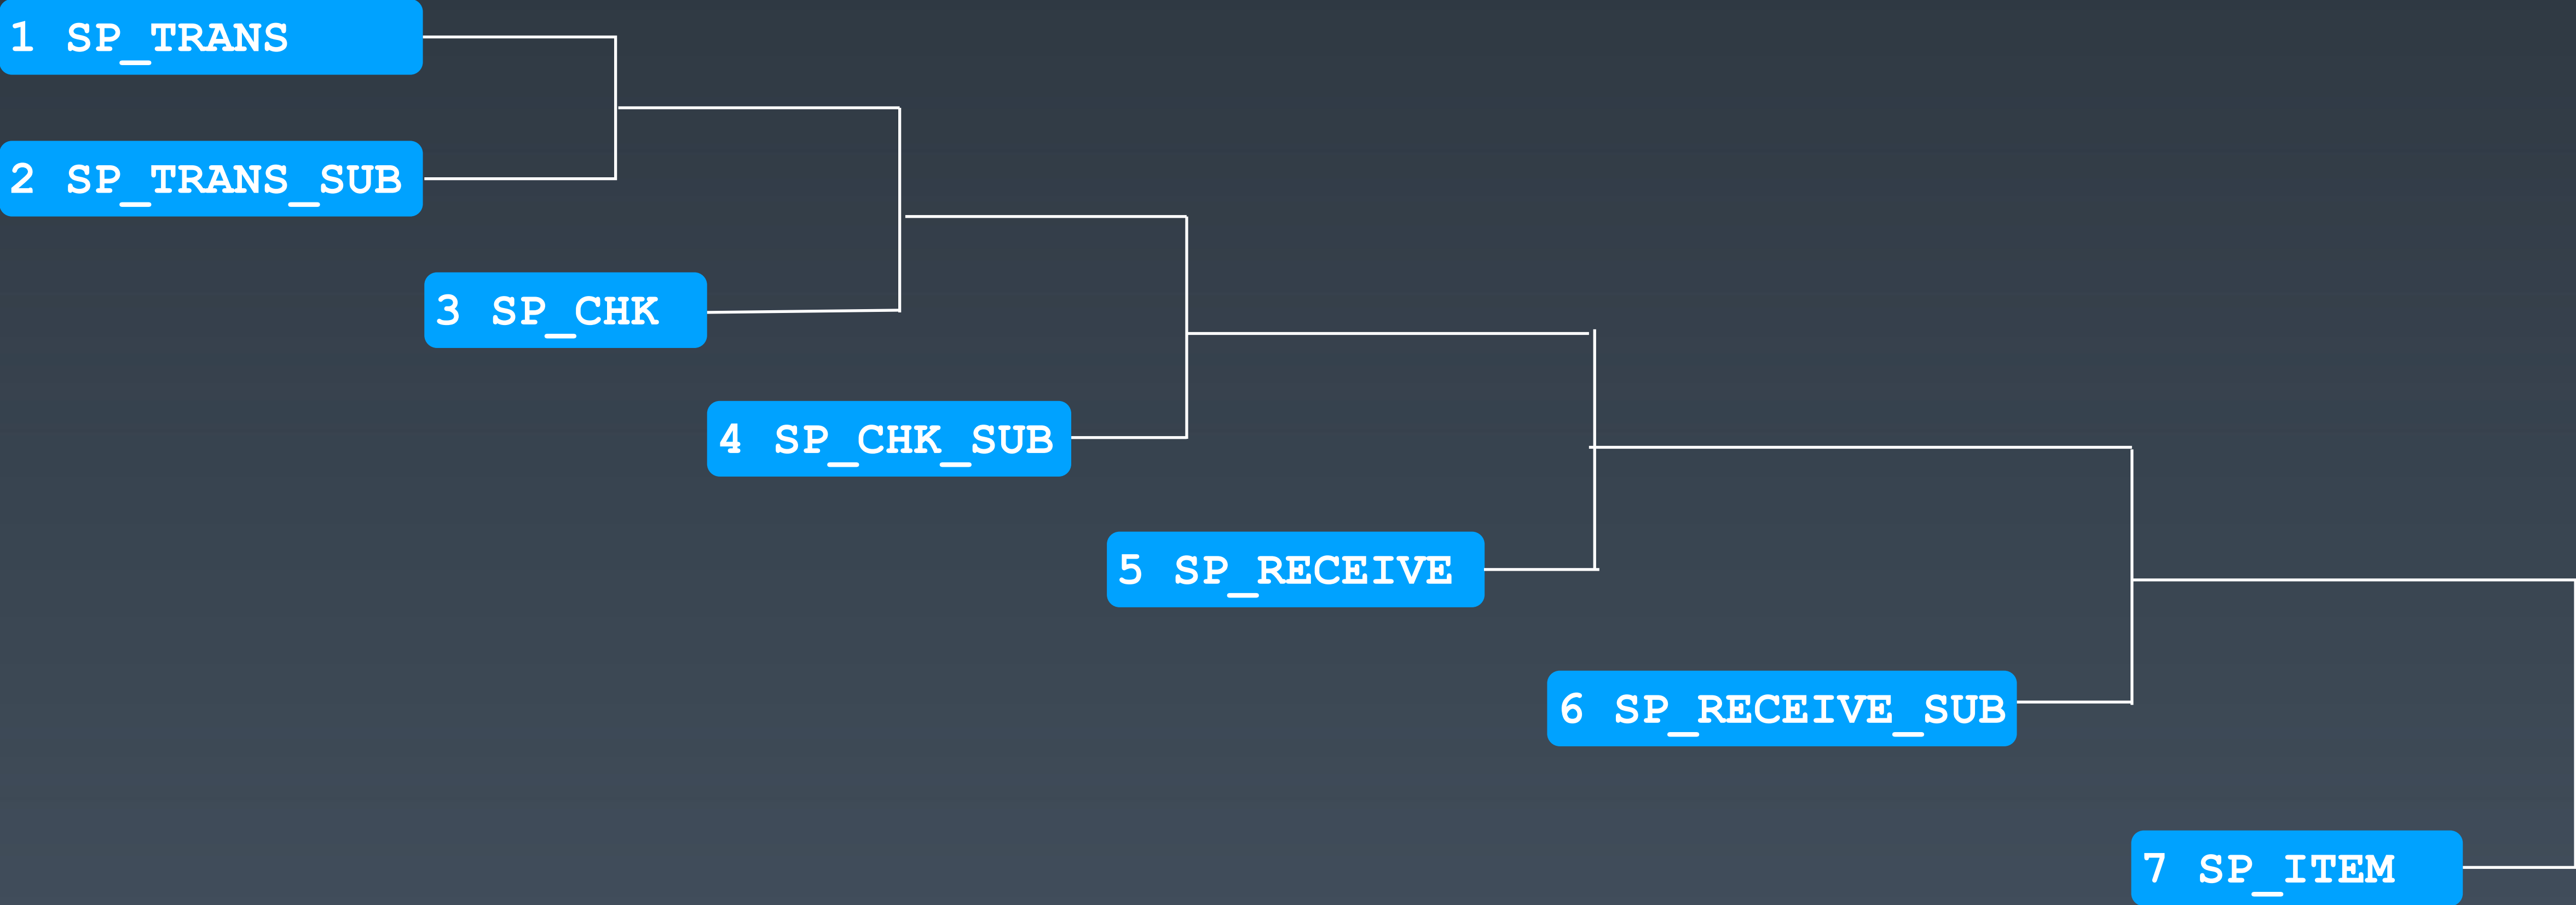
4 SP_CHK_SUB

5 SP_RECEIVE

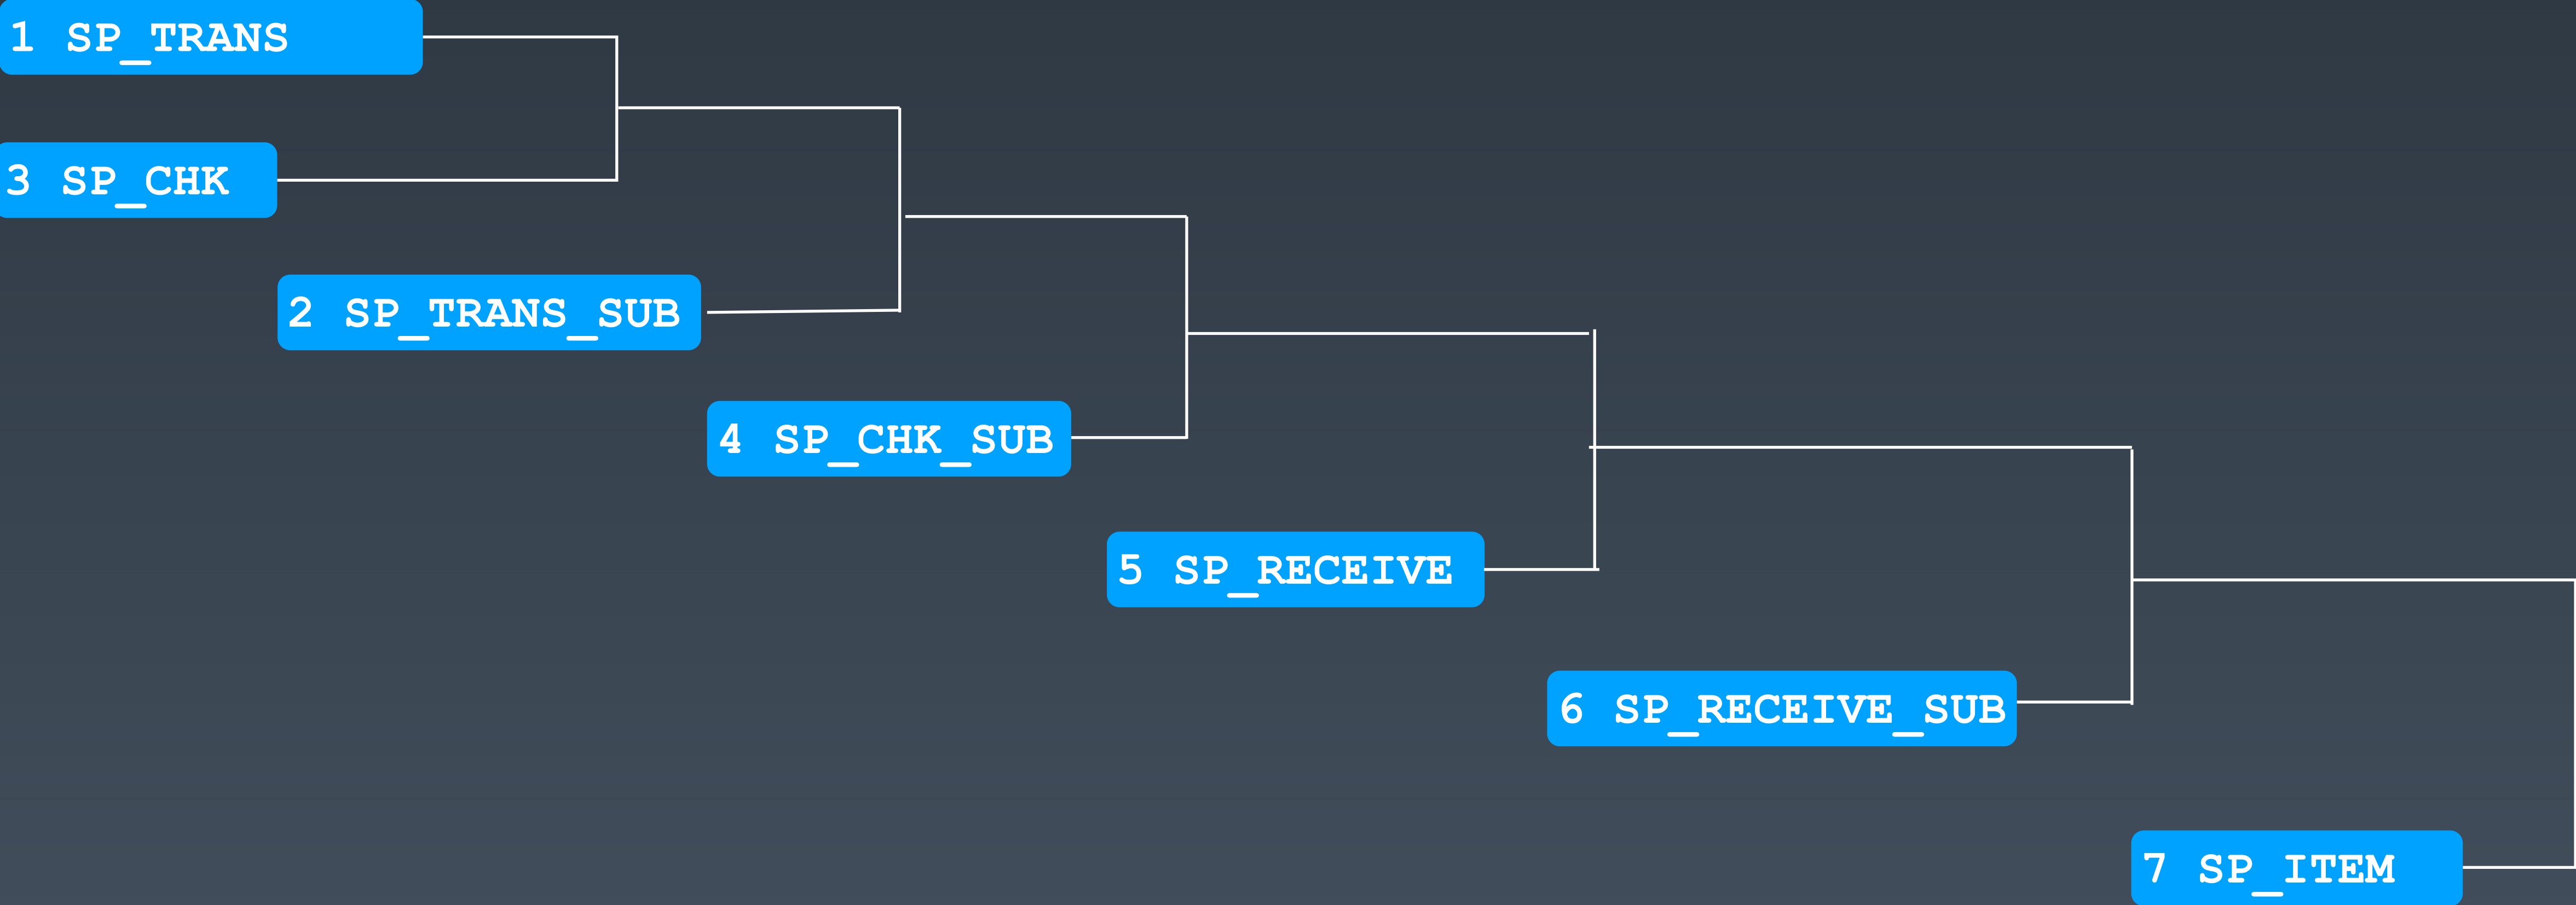
6 SP_RECEIVE_SUB

7 SP_ITEM

SQL背后的世界：多表如何联合



SQL背后的世界：多表如何联合



SQL背后的世界：理解优化器

1 SP_TRANS

2 SP_TRANS_SUB

3 SP_CHK


4 SP_CHK_SUB

5 SP_RECEIVE

6 SP_RECEIVE_SUB

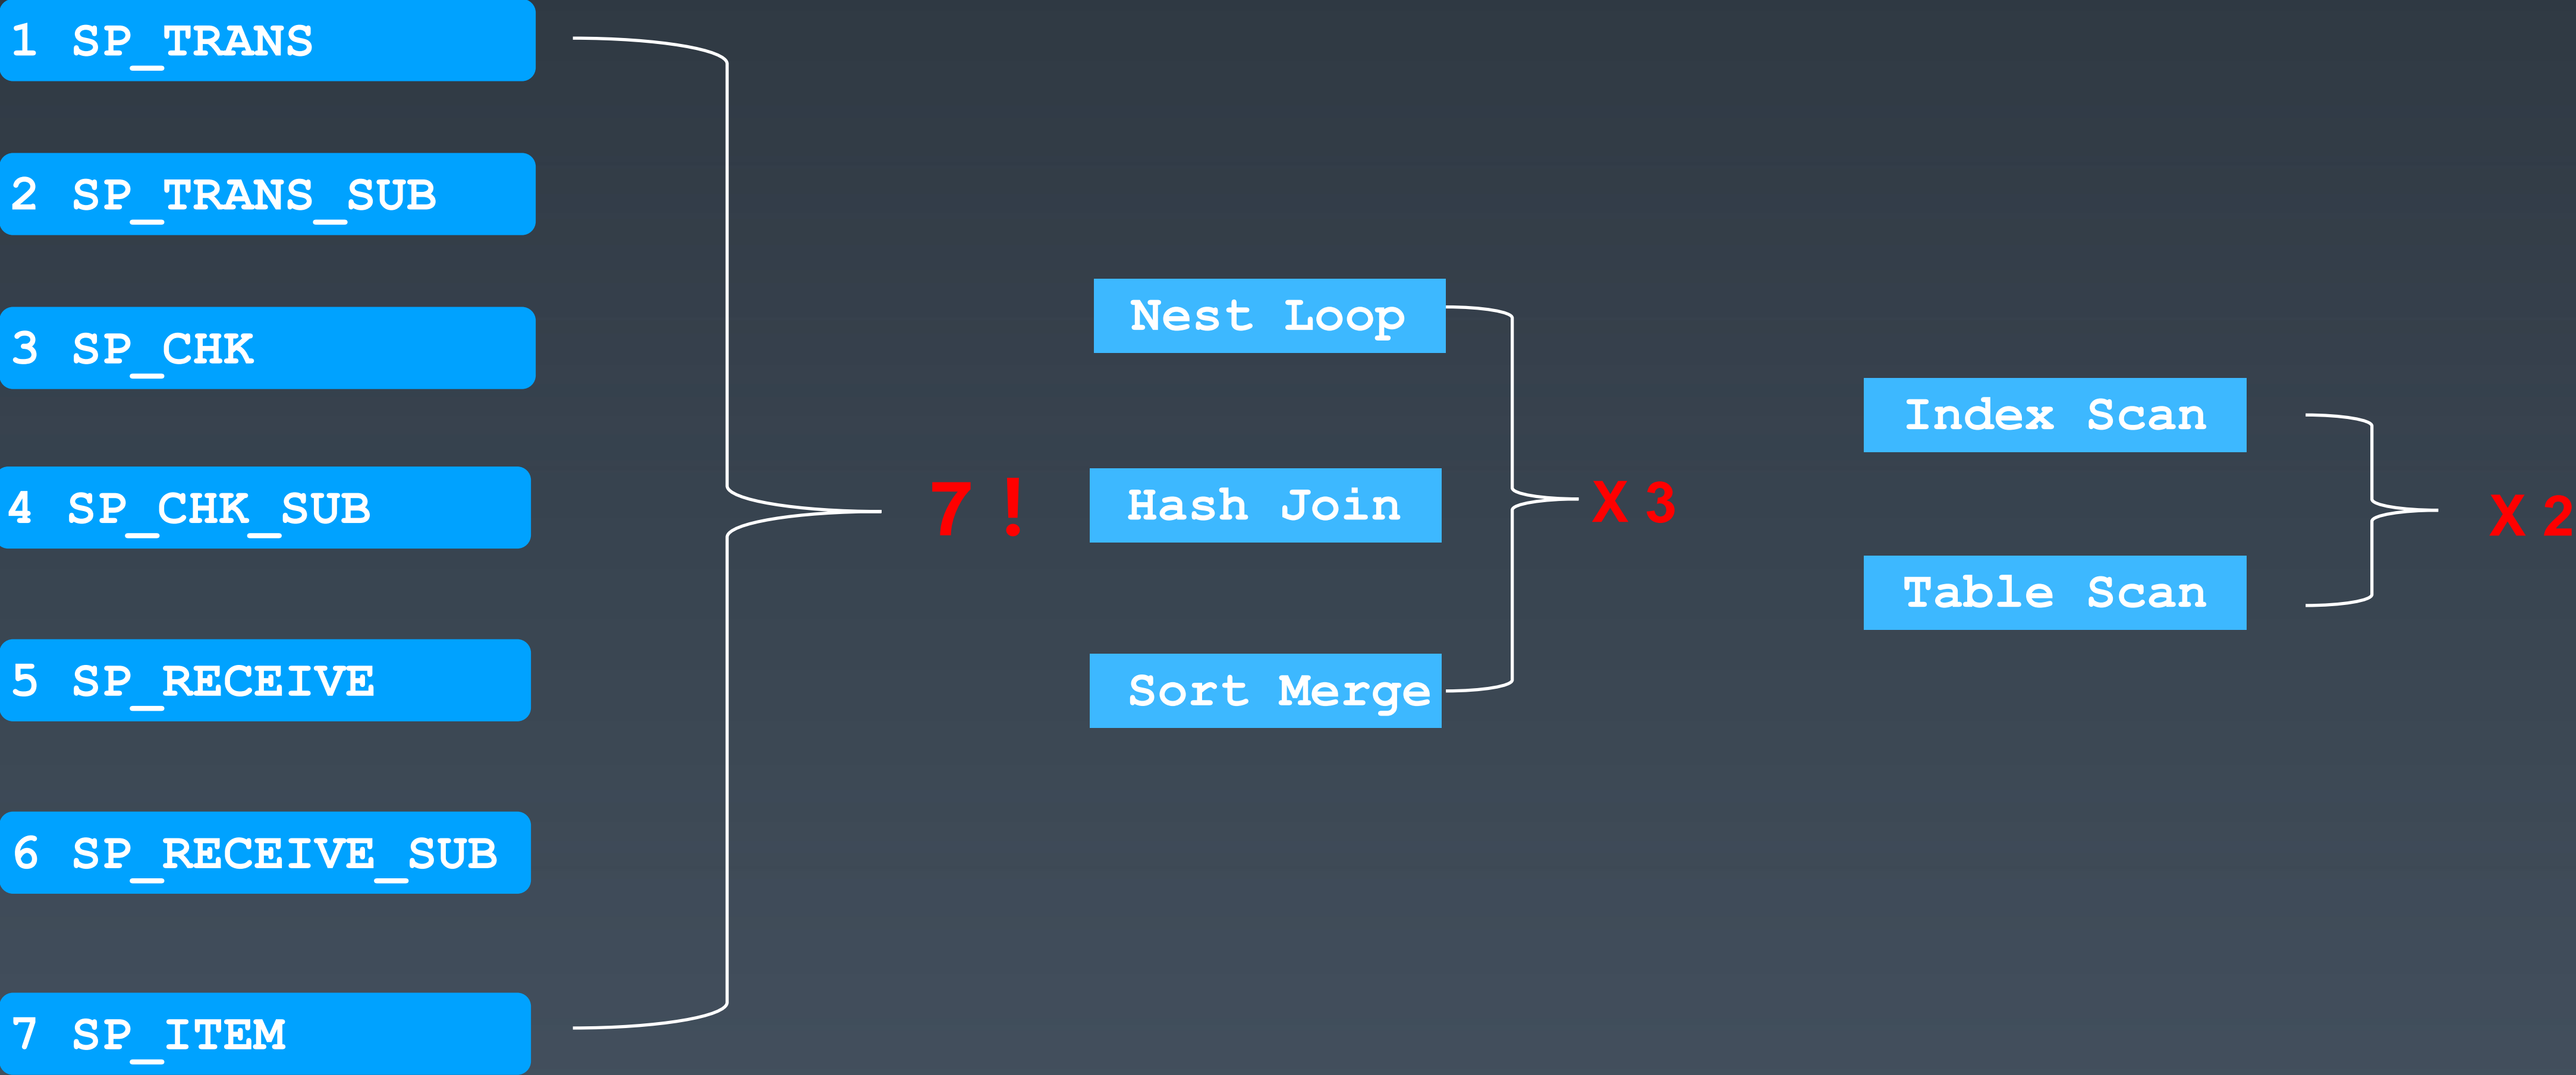
7 SP_ITEM

7 !



_optimizer_max_permutations

SQL背后的世界：理解优化器



SQL背后的世界：理解优化器

Join order[88]: SP_RECEIVE [SP_RECEIVE] SP_RECEIVE_SUB [SP_RECEIVE_SUB] SP_CHK [SP_CHK] SP_CHK_SUB
[SP_CHK_SUB] SP_TRANS [SP_TRANS] SP_TRANS_SUB [SP_TRANS_SUB] SP_ITEM [SP_ITEM]

Now joining: SP_CHK_SUB [SP_CHK_SUB] *****

NL Join

Outer table: cost: 1863 **cdn**: 18225 rcz: 124 resp: 1863
Inner table: SP_CHK_SUB Access path: tsc Resc: 60 Join resc: 1095363 Resp: 1095363
Join cardinality: 1 = outer (18225) * inner (36532) * sel (1.3146e-015) [flag=0]
Best NL **cost**: 1095363 resp: 1095363

SM Join



Outer table: resc: 1863 cdn: 18225 rcz: 124 deg: 1 resp: 1863
Inner table: SP_CHK_SUB resc: 60 cdn: 36532 rcz: 48 deg: 1 resp: 60
Merge join Cost: 3856 Resp: 3856

HA Join

Outer table: resc: 1863 cdn: 18225 rcz: 124 deg: 1 resp: 1863
Inner table: SP_CHK_SUB resc: 60 cdn: 36532 rcz: 48 deg: 1 resp: 60
Hash join one ptn: 3414 Deg: 1 (sides swapped)

hash_area: 16 buildfrag: 268 probefrag: 303 ppasses: 17
Hash join Resc: 5337 Resp: 5337

数据背后：藏身后台的索引和性能



SQL Details: cd3sb60yu30ab

Switch to SQL ID

Text

```
INSERT INTO "EDIDS" VALUES (:A0 ,:A1 ,:A2 ,:A3 ,:A4 ,:A5 ,:A6 ,:A7 ,:A8 ,:A9 ,:A10 ,:A11 ,
```

Details

Select the plan hash value to see the details below. Plan Hash Value There are multiple plans found for this SQL statement.

Statistics **Activity** **Plan** **Plan Control** **Tuning History** **SQL Monitoring**

Summary

Top SQL: Concurrency

Actions

Select All | Select None

| Select Activity (%) | SQL ID | SQL Type |
|--------------------------------|---------------|----------|
| <input type="checkbox"/> 95.64 | cd3sb60yu30ab | INSERT |
| <input type="checkbox"/> 1.65 | b75zqwgfg6hms | INSERT |
| <input type="checkbox"/> .24 | b75zqwgfg6hms | INSERT |
| <input type="checkbox"/> .21 | cd3sb60yu30ab | INSERT |
| <input type="checkbox"/> .14 | b75zqwgfg6hms | INSERT |

Top Sessions: Concurrency

View Top Sessions

| Activity (%) | Session ID | User Name | Program |
|--------------|------------|-------------------------------|---------|
| .38 6095 | SAPSR3 | disp+work@gmpap05 (TNS V1-V3) | |
| .38 8458 | SAPSR3 | disp+work@gmpap03 (TNS V1-V3) | |
| .37 3960 | SAPSR3 | disp+work@gmpap01 (TNS V1-V3) | |
| .37 8146 | SAPSR3 | disp+work@gmpap03 (TNS V1-V3) | |
| .37 9883 | SAPSR3 | disp+work@gmpap05 (TNS V1-V3) | |
| .37 2535 | SAPSR3 | disp+work@gmpap05 (TNS V1-V3) | |
| .37 8204 | SAPSR3 | disp+work@gmpap01 (TNS V1-V3) | |

索引分裂：高并发时数据库的表征

高事务并发的典型案例

- 行锁竞争 row lock contention;
- 索引分裂 index contention;
- 前者源自应用，后者源于索引;

Buffer Busy Waits

- 索引是主要矛盾;

Top 10 Foreground Events by Total Wait Time

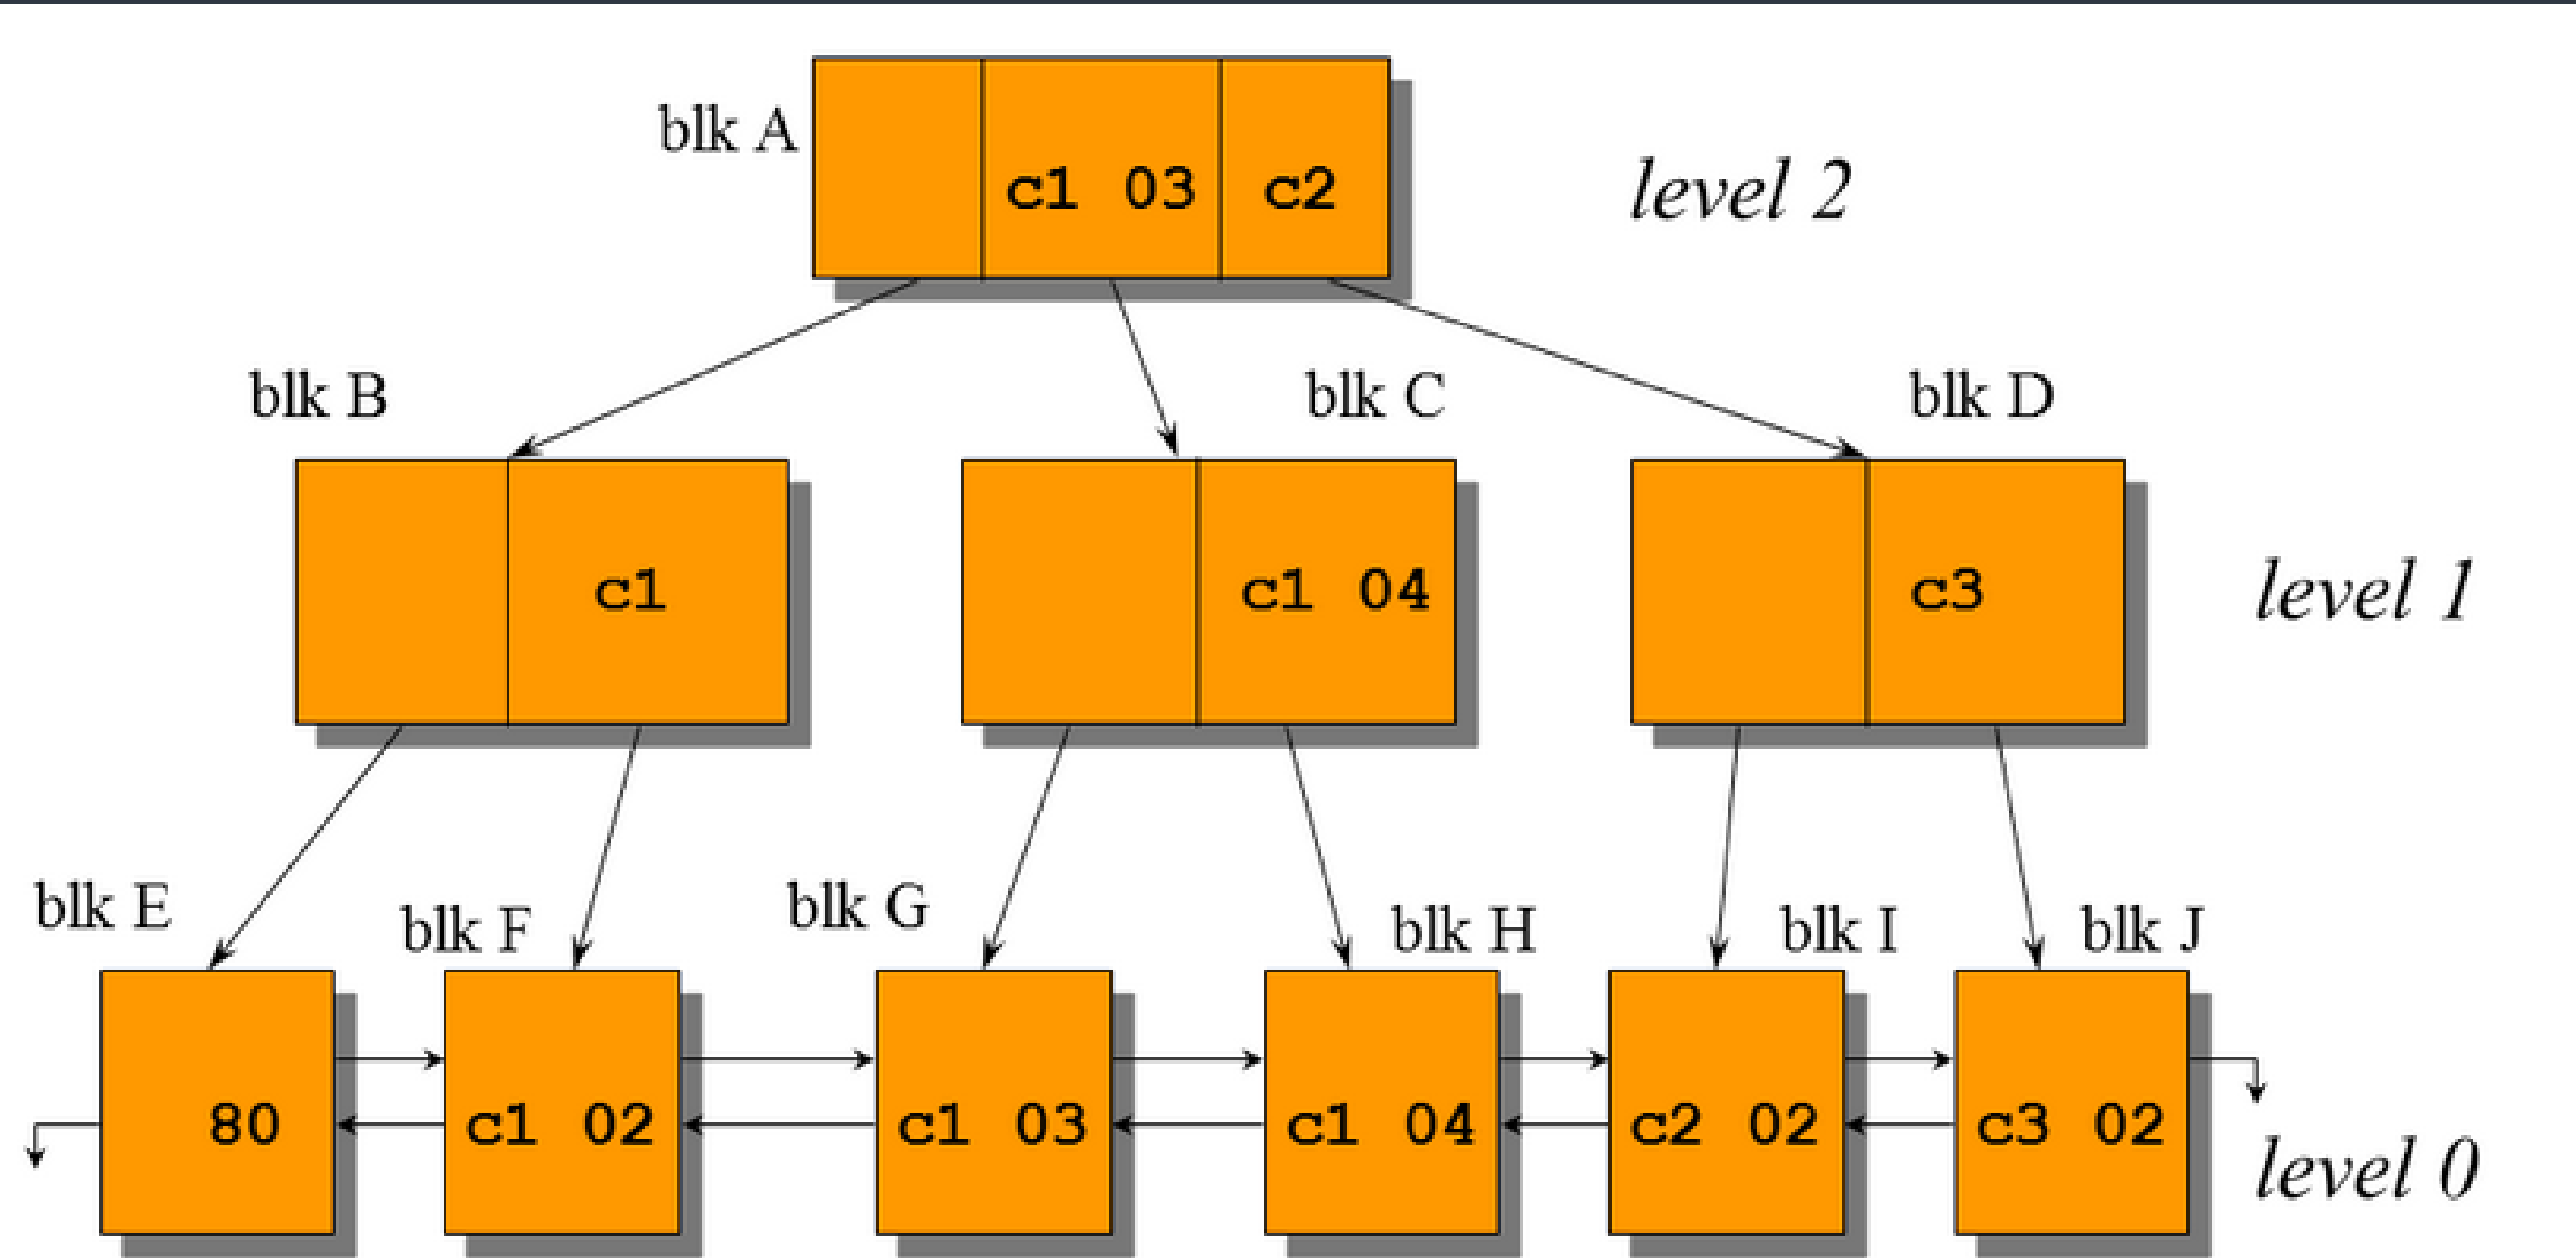
| Event | Waits | Total Wait Time (sec) | Wait Avg(ms) | % DB time | Wait Class |
|-------------------------------|-----------|-----------------------|--------------|-----------|---------------|
| DB CPU | | 28.1K | | 24.3 | |
| enq: TX - row lock contention | 40,705 | 24.1K | 592 | 20.8 | Application |
| buffer busy waits | 113,910 | 11.4K | 100 | 9.9 | Concurrency |
| enq: TX - index contention | 60,697 | 7928.7 | 131 | 6.9 | Concurrency |
| log file sync | 745,655 | 5530.8 | 7 | 4.8 | Commit |
| db file sequential read | 2,074,730 | 5518.3 | 3 | 4.8 | User I/O |
| enq: HW - contention | 430 | 4407.3 | 10250 | 3.8 | Configuration |
| latch: cache buffers chains | 61,510 | 2593.1 | 42 | 2.2 | Concurrency |

Segments by Buffer Busy Waits

- % of Capture shows % of Buffer Busy Waits for each top segment compared
- with total Buffer Busy Waits for all segments captured by the Snapshot

| Owner | Tablespace Name | Object Name | Subobject Name | Obj. Type | Buffer Busy Waits | % of Capture |
|----------|-----------------|-----------------------|-------------------------|-----------------|-------------------|--------------|
| P2BEMADM | BEM_DEFUSER_DAT | IDX_PRODUCT_04 | | INDEX | 42,485 | 36.00 |
| P2BEMADM | BEM_PRDHIST_DAT | IDX_PRODUCTHISTORY_01 | PRODUCTHISTORY_20180427 | INDEX PARTITION | 22,478 | 19.05 |
| P2BEMADM | BEM_DEFUSER_DAT | IDX_LOT_08 | | INDEX | 17,420 | 14.76 |
| P2BEMADM | BEM_CUSTOMS_IDX | CT_MESSAGELOG_PK | P_20180427 | INDEX PARTITION | 12,010 | 10.18 |
| P2BEMADM | BEM_LOTHIST_DAT | IDX_LOTHISTORY_02 | LOTHISTORY_20180427 | INDEX PARTITION | 9,865 | 8.36 |

索引分裂：高并发时数据库的表征



Segments by Buffer Busy Waits

- % of Capture shows % of Buffer Busy Waits for each top segment compared
- with total Buffer Busy Waits for all segments captured by the Snapshot

| Owner | Tablespace Name | Object Name | Subobject Name | Obj. Type | Buffer Busy Waits | % of Capture |
|----------|-----------------|-----------------------|-------------------------|-----------------|-------------------|--------------|
| P2BEMADM | BEM_DEFUSER_DAT | IDX_PRODUCT_04 | | INDEX | 42,485 | 36.00 |
| P2BEMADM | BEM_PRDHIST_DAT | IDX_PRODUCTHISTORY_01 | PRODUCTHISTORY_20180427 | INDEX PARTITION | 22,478 | 19.05 |
| P2BEMADM | BEM_DEFUSER_DAT | IDX_LOT_08 | | INDEX | 17,420 | 14.76 |
| P2BEMADM | BEM_CUSTOMS_IDX | CT_MESSAGELOG_PK | P_20180427 | INDEX PARTITION | 12,010 | 10.18 |
| P2BEMADM | BEM_LOTHIST_DAT | IDX_LOTHISTORY_02 | LOTHISTORY_20180427 | INDEX PARTITION | 9,865 | 8.36 |

- 常规表的单列索引
 - IDX_PRODUCT_04 – product -LASTTIMEKEY;
 - IDX_LOT_08 – lot - LASTTIMEKEY;
- 时间范围分区表的复合索引
 - IDX_PRODUCTHISTORY_01 -producthistory - (timekey,productname) ;
 - IDX_LOTHISTORY_02 - lothistory- (timekey,eventname,machinename);
- LASTTIMEKEY/TIMEKEY是时间戳，单调递增;
- Index contention 是索引分裂同时有其他会话尝试更新。最容易产生等待的是单调递增的索引，因为每次插入都在索引的最右边。

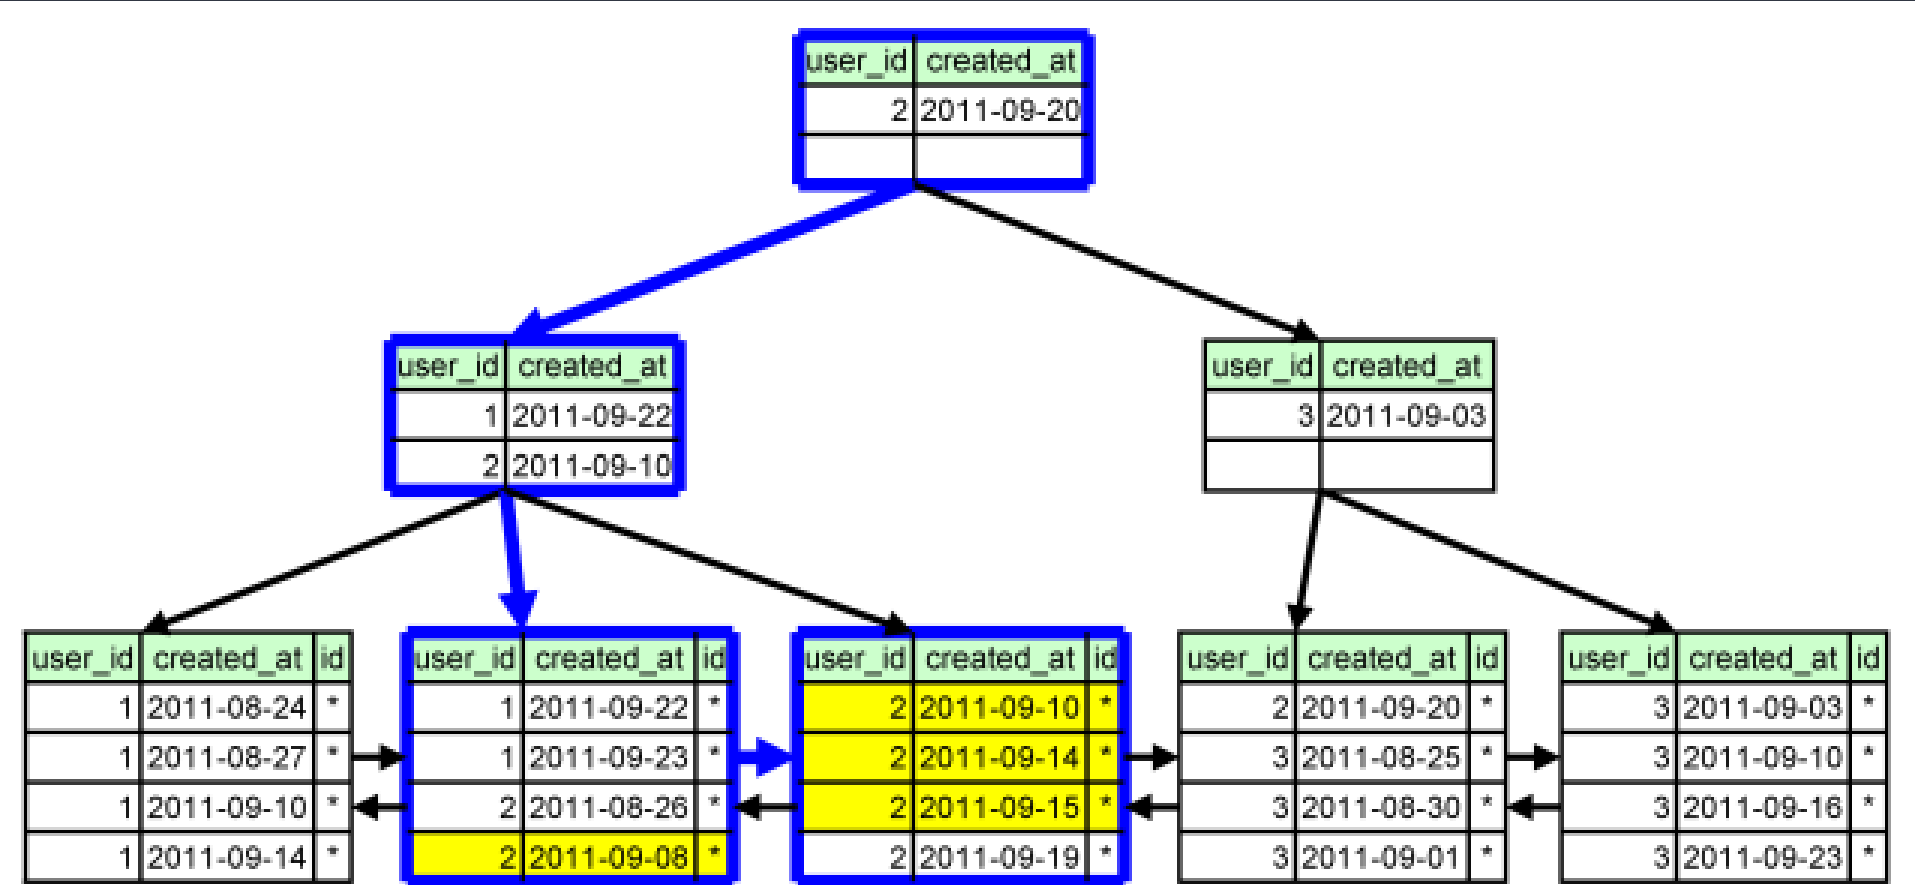
解决方案：追根溯源分解竞争

- 数据库级解决方案：降低 buffer busy wait / index contention，也就是打散热点索引，把相近 timekey 对应的索引块分散；
- IDX_PRODUCT_04 / IDX_LOT_08，通过建立hash partition index，就可以把索引块分布到不同的分区上，大幅降低争用；
 - Create index IDX_PRODUCT_04 on PRODUCT(lasttimekey) global partition by hash(lasttimekey) partitions 16;
 - Create index IDX_LOT_08 on LOT(lasttimekey) global partition by hash(lasttimekey) partitions 16;
- Producthistory / lothistory 是分区表，无法直接创建hash子分区索引，除非把整个表重建为 range-hash的复合分区，但是这么做改动太大；
- 分析SQL，发现producthistory表的访问基本都是两个栏位productname和timekey的，因此把索引重建为 productname 在前，timekey 在后就可解决问题。
 - Create index idx_producthistory_01(productname,timekey);

解决方案：追根溯源分解竞争

• LOTHISTORY 表的索引优化

- SQL主要以 timekey 为条件，很多不包括eventname, machinename;
- 必须有合适的索引可以应用 timekey 条件;
- 可以增加前缀列，让SQL走index skip scan;
- 但是前缀列的唯一值不能太大, 否则index skip scan的额外成本高;
- 前缀列的唯一值也不能太小，否则起不到分散索引块的作用。
- 考虑使用8-32之间的值。



Top 10 Foreground Events by Total Wait Time

| Event | Waits | Total Wait Time (sec) | Wait Avg(ms) | % DB time | Wait Class |
|-------------------------------|-----------|-----------------------|--------------|-----------|-------------|
| DB CPU | | 16.4K | | 86.4 | |
| enq: TX - row lock contention | 25,638 | 1725 | 67 | 9.1 | Application |
| SQL*Net more data from client | 3,042,850 | 660 | 0 | 3.5 | Network |
| log file sync | 668,232 | 239.5 | 0 | 1.3 | Commit |
| db file sequential read | 429,029 | 160.6 | 0 | .8 | User I/O |
| SQL*Net message from dblink | 146,166 | 128.9 | 1 | .7 | Network |
| gc current grant busy | 560,576 | 96.7 | 0 | .5 | Cluster |

索引分裂：深入数据库的原理

- 右向增长索引

经常在索引分裂时引起严重竞争；

尤其是以序列单调递增的方式；

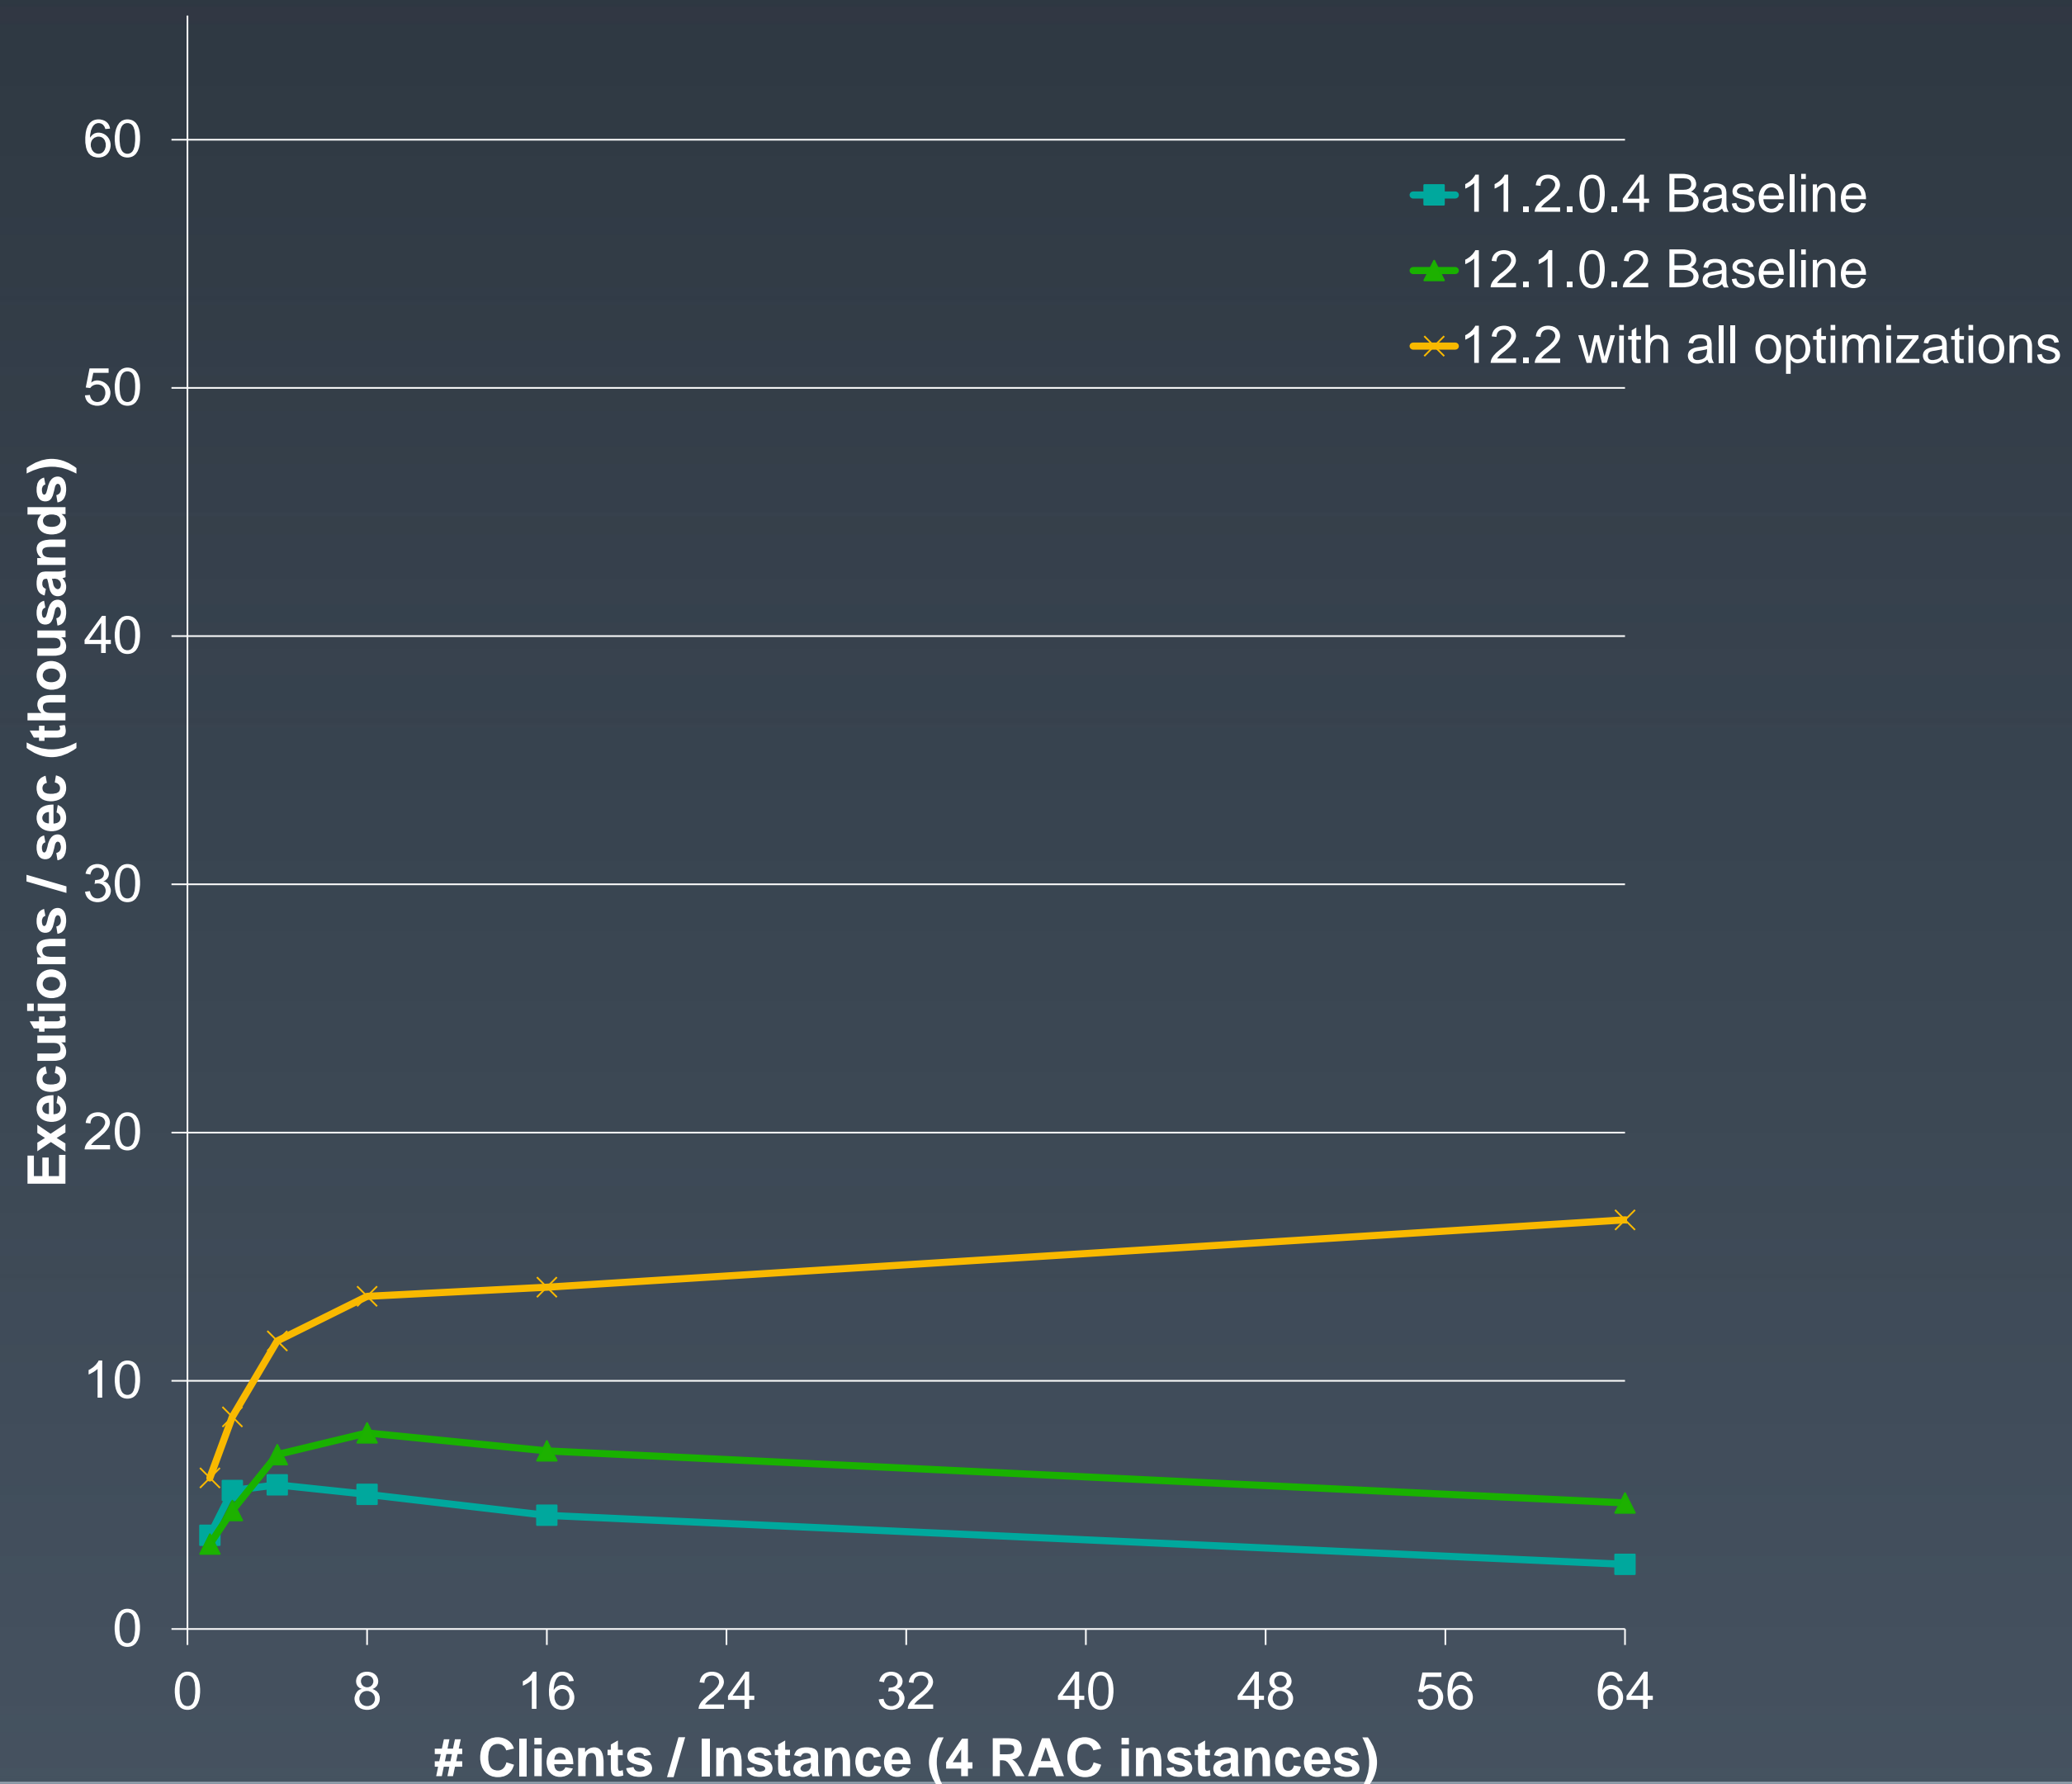
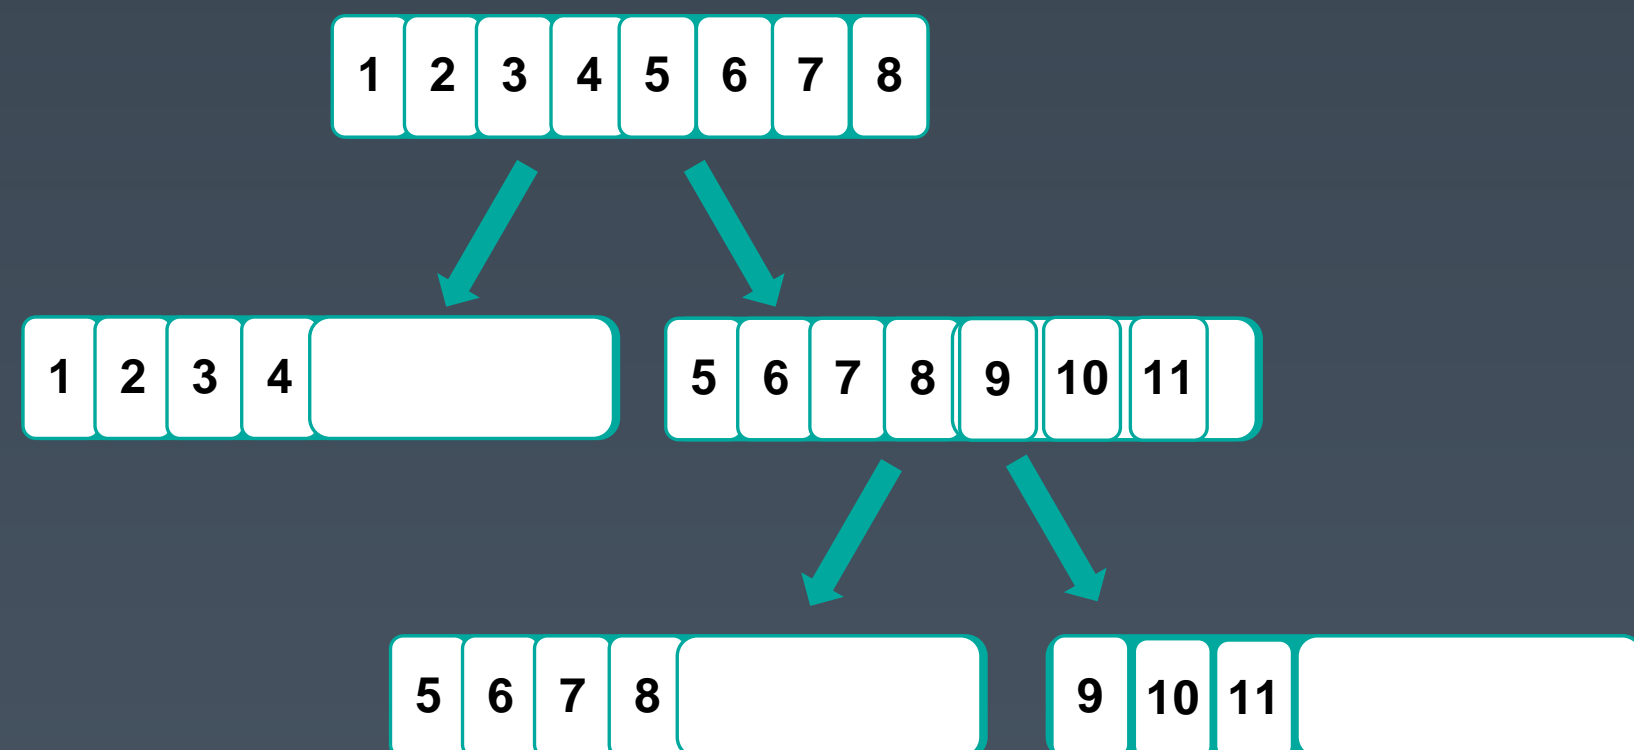
- 优化了叶块分裂算法

缩减集群消息传递；

- 使用更佳的索引特性

18c 的 Scalable Sequences；

模拟 Scalable 方式创建散列索引；



特性增强：可扩展序列分散索引竞争

在Scaleable Sequence 中指定SCALE时，在传统的序列前增加了6位数字

- 前3位是由RAC里的实例号产生
- 随后3位由Session的SID产生


```
SQL> create sequence seq_eygle  
      start with 1 increment by 1 scale;
```

```
SQL>select seq_eygle.nextval from dual;
NEXTVAL
```

[illegible]

```
SQL>select instance_number from v$instance;
```

```
SQL>select sid from v$mystat where rownum=1;
```



```
SQL>select sid from v$mystat where rownum=1;
```

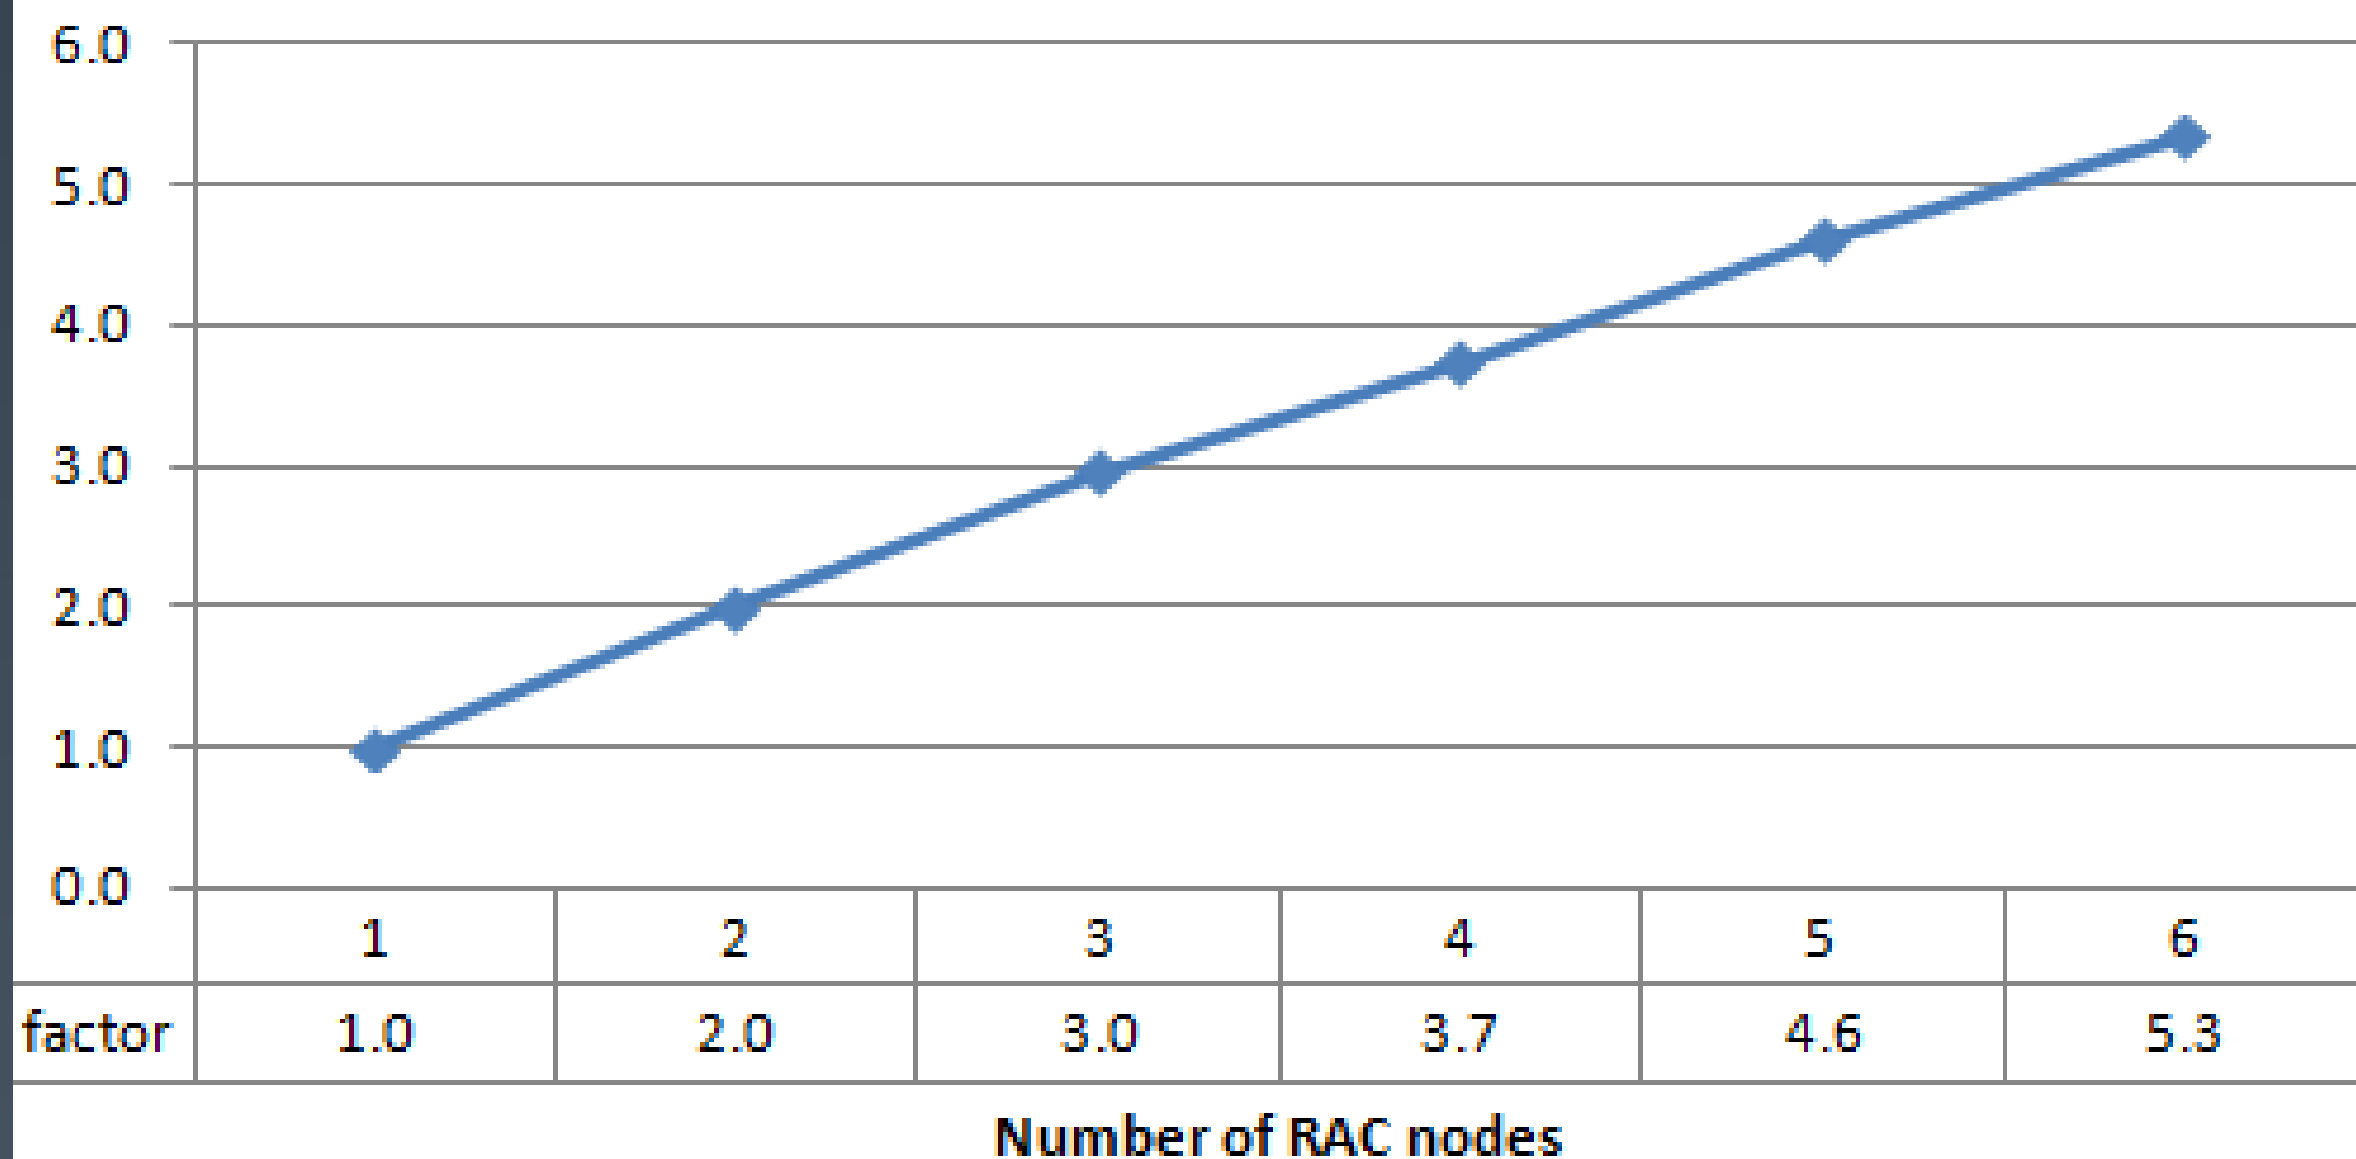
```
SQL>select seq_eygle.nextval from dual;
NEXTVAL
```

101264000000000000000000000005

特性增强：Sharded RAC分散集群竞争

- 将分片能力引入到RAC集群实例中
 - 指定了分区键值的SQL查询将被路由到特定的实例;
 - 分区可以避免竞争减少跨实例的访问;
- 不包含分区键值的请求会透明的被处理
- 以最小的应用改变提供分布式性能

Figure 4 - RAC DB Performance with RAC Affinity



Instance 1
Partition P1

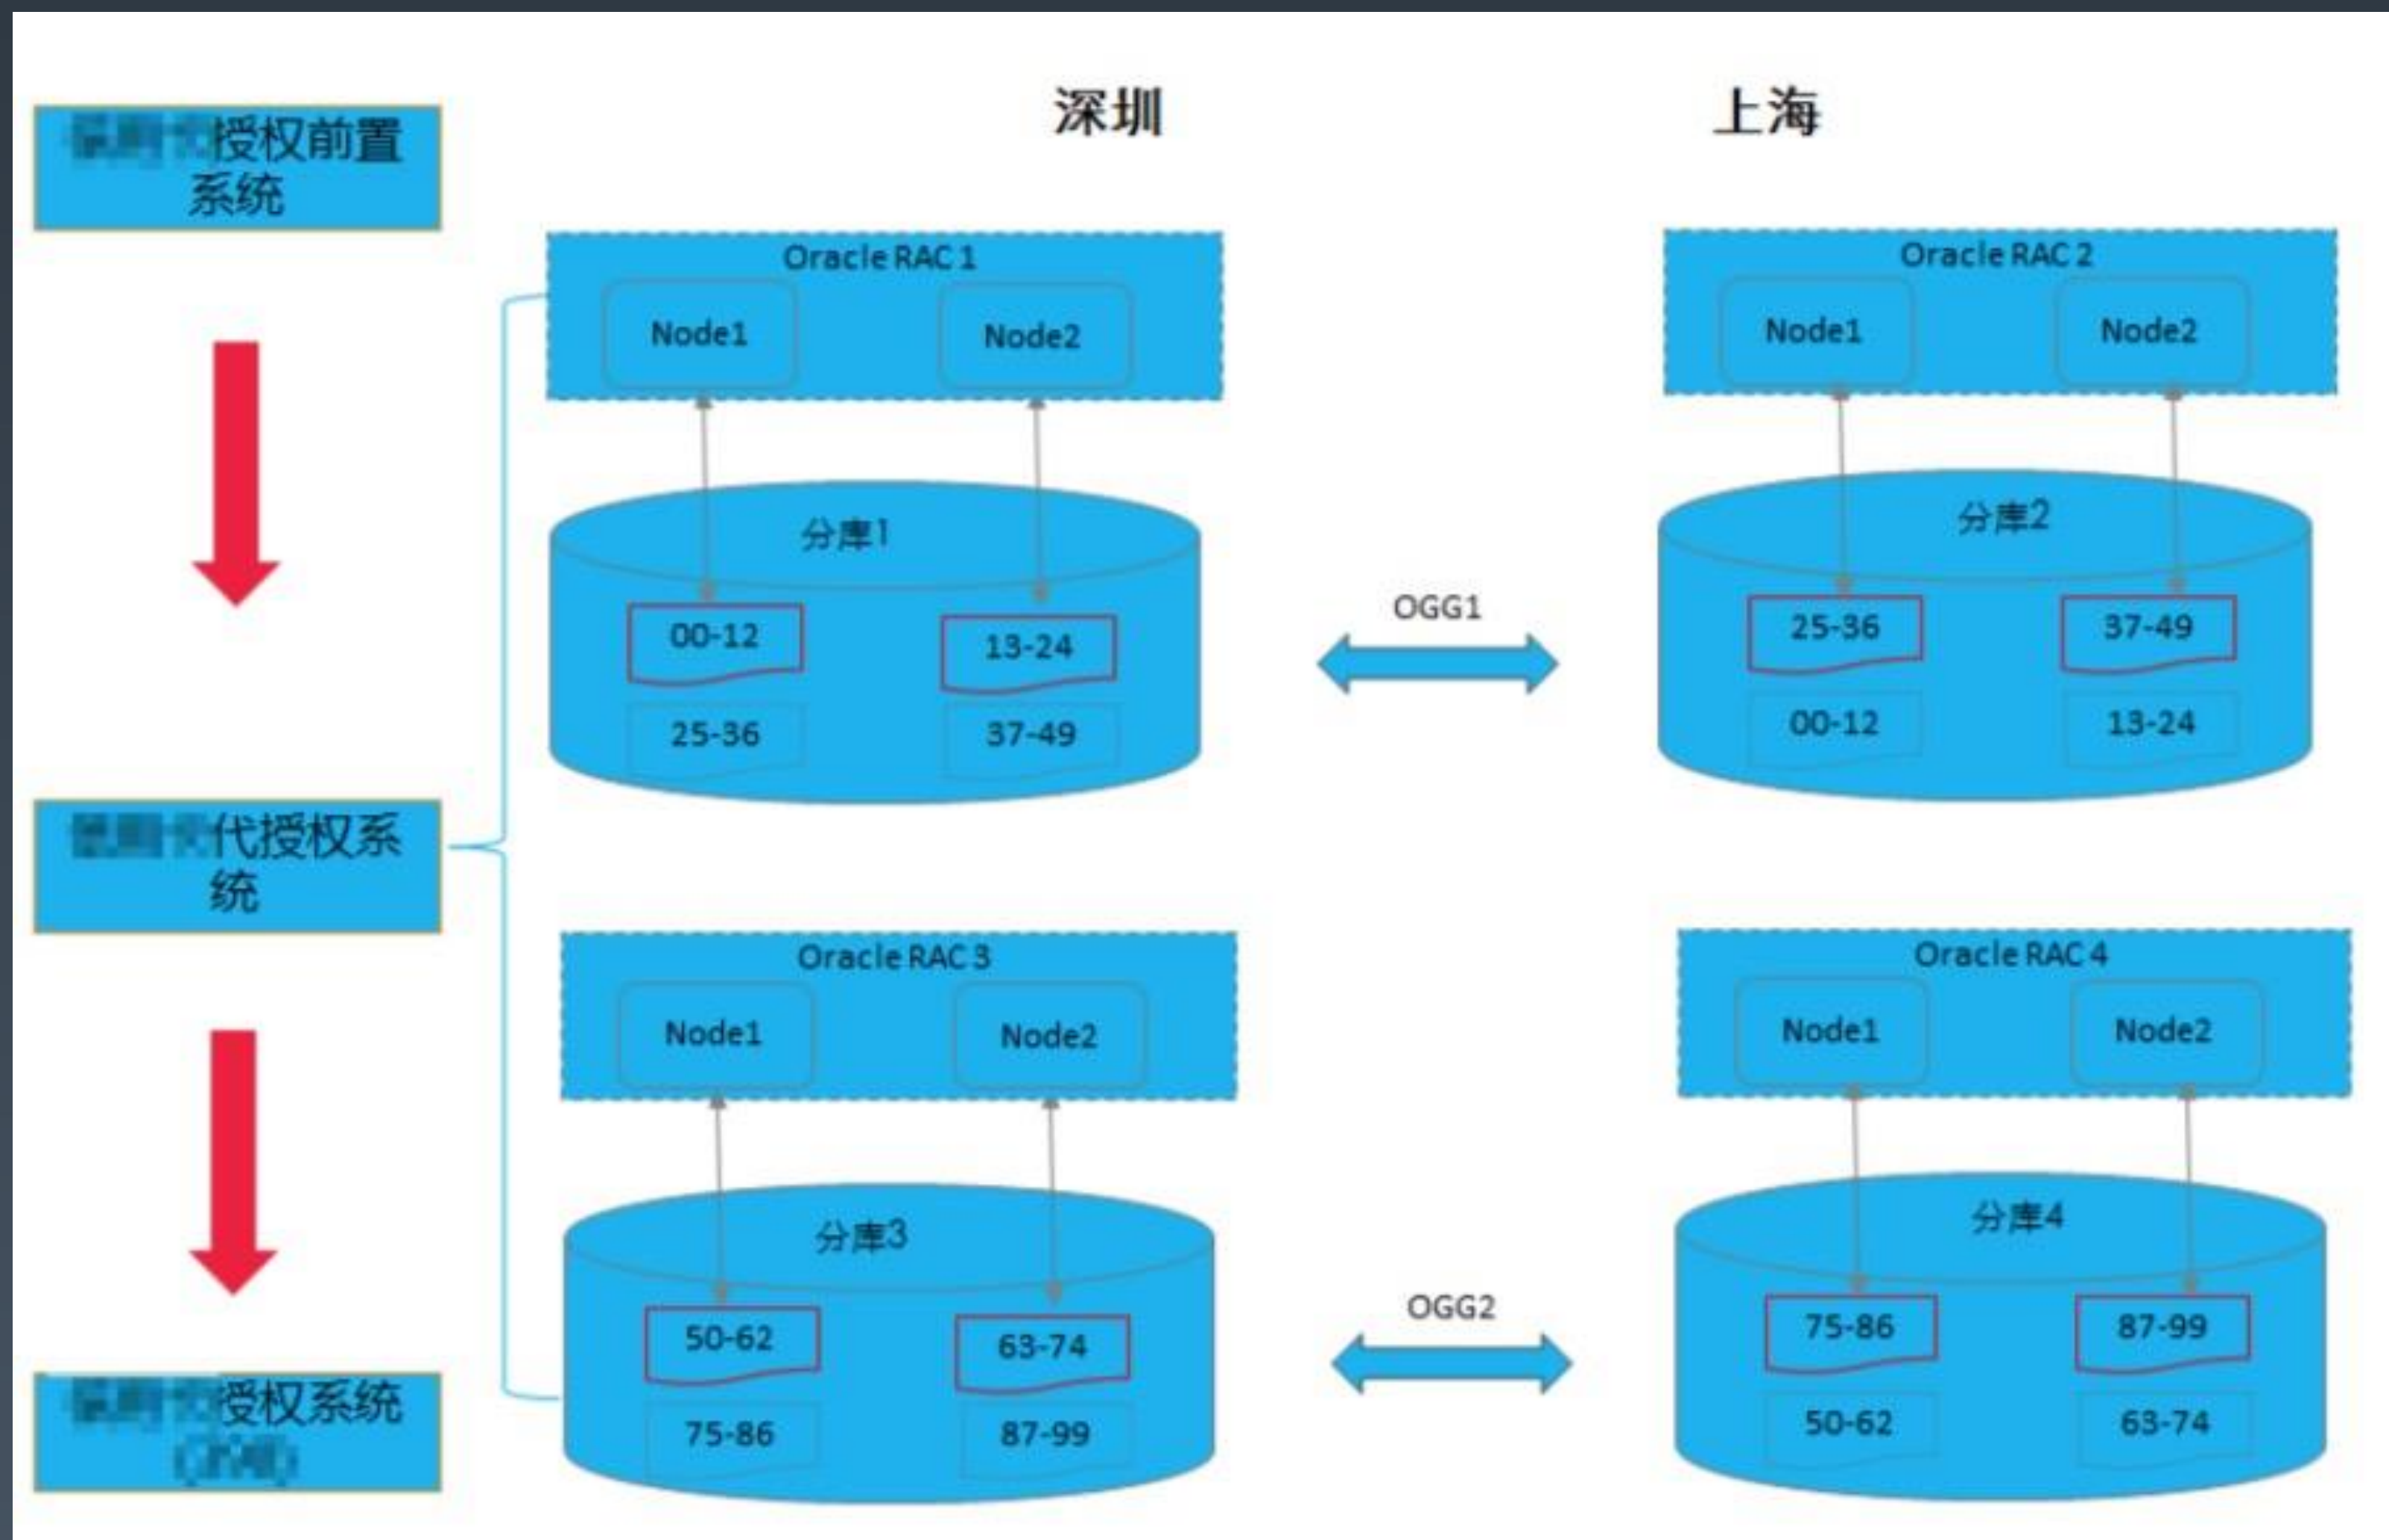
Instance 2
Partition P2

Instance 3
Partition P3



Oracle RAC Database

设计增强：通过分库分表获得线性性能



- 正常情况下
 - 每套RAC双节点;
 - 2套RAC组成双活架构;
 - 通过应用使RAC的各节点访问的数据不交叉;
 - 应用根据客户号后2位分割路由数据;
- 单节点失效
 - 同一RAC中的其他节点接管;
- 站点失效;
 - 切换路由到双活的异地节点;

目 录

- Oracle 的变革之路
- 性能优化与数据库进化
- **Oracle 19c新特性**

数据库上的明珠：SQL 的优化器

SQL

AUTOMATION

备库的自动DML转发

- 使得主备环境的作用进一步增强

失控SQL的执行计划自动隔离

- 实现对于性能的主动防御

自动化的执行计划管理 - Automated Plan Management

- 自动捕获所有可重用SQL的执行计划
- 查找可选计划选择最佳计划

实时的统计信息 - Real Time Statistics

- 近实时的统计信息刷新

SQL ENHANCEMENTS

使用统计信息支持查询

- 使用已经收集的统计信息，避免扫描大量数据；例如，`select count(*) from emp`
- 可以极大提升某些检查查询的性能，Min, Max, Count, Approximate_Count_distinct

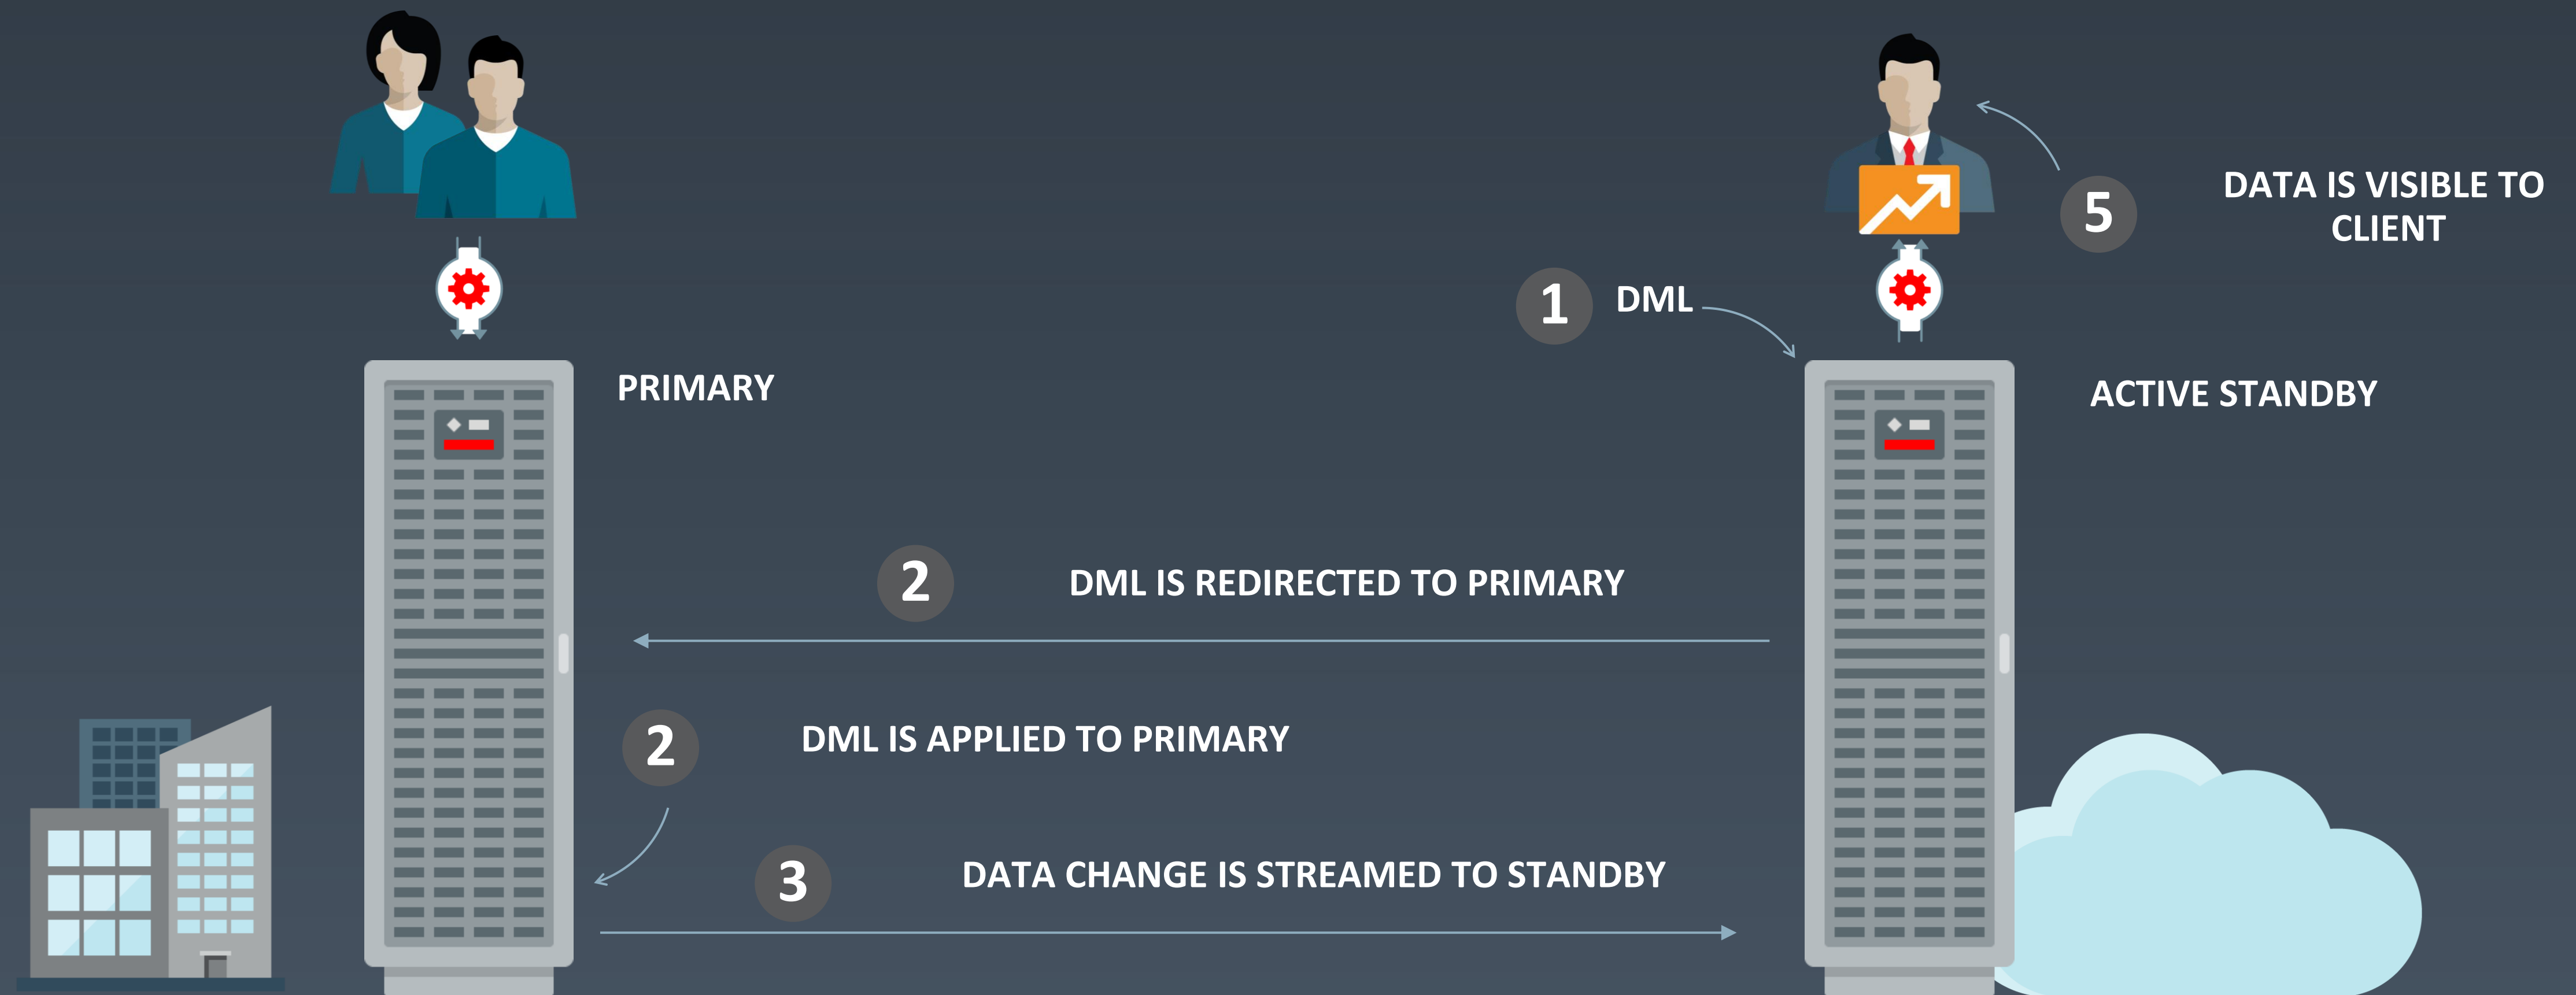
混合表的支持

- 通过分区指向外部存储对象，实现外部数据访问

负载分担：在备库透明的支持DML语句

NEW IN
19^C

- 在 ADG 上发出的 DML 重定向到主库执行，备库等待主库日志传递到备库并应用



负载分担：在备库透明的支持DML语句

NEW IN
19^c

```
SQL> select * from enmotech;  
ID NAME
```

```
1 EYGLE  
2 KAMUS  
3 Yangtingkun
```

```
SQL> insert into enmotech values (4, 'ORA-600');  
ORA-16000: database or pluggable database open  
for read-only access
```

```
SQL> alter session set adg_redirect_dml=true;  
Session altered.
```

```
SQL> set timing on  
SQL> insert into enmotech values (4, 'ORA-600');  
1 row created.
```

Elapsed: **00:00:01.21**

```
SQL> select * from enmotech;  
ID NAME
```

```
4 ORA-600  
1 EYGLE  
2 KAMUS  
3 Yangtingkun
```

Elapsed: 00:00:00.01

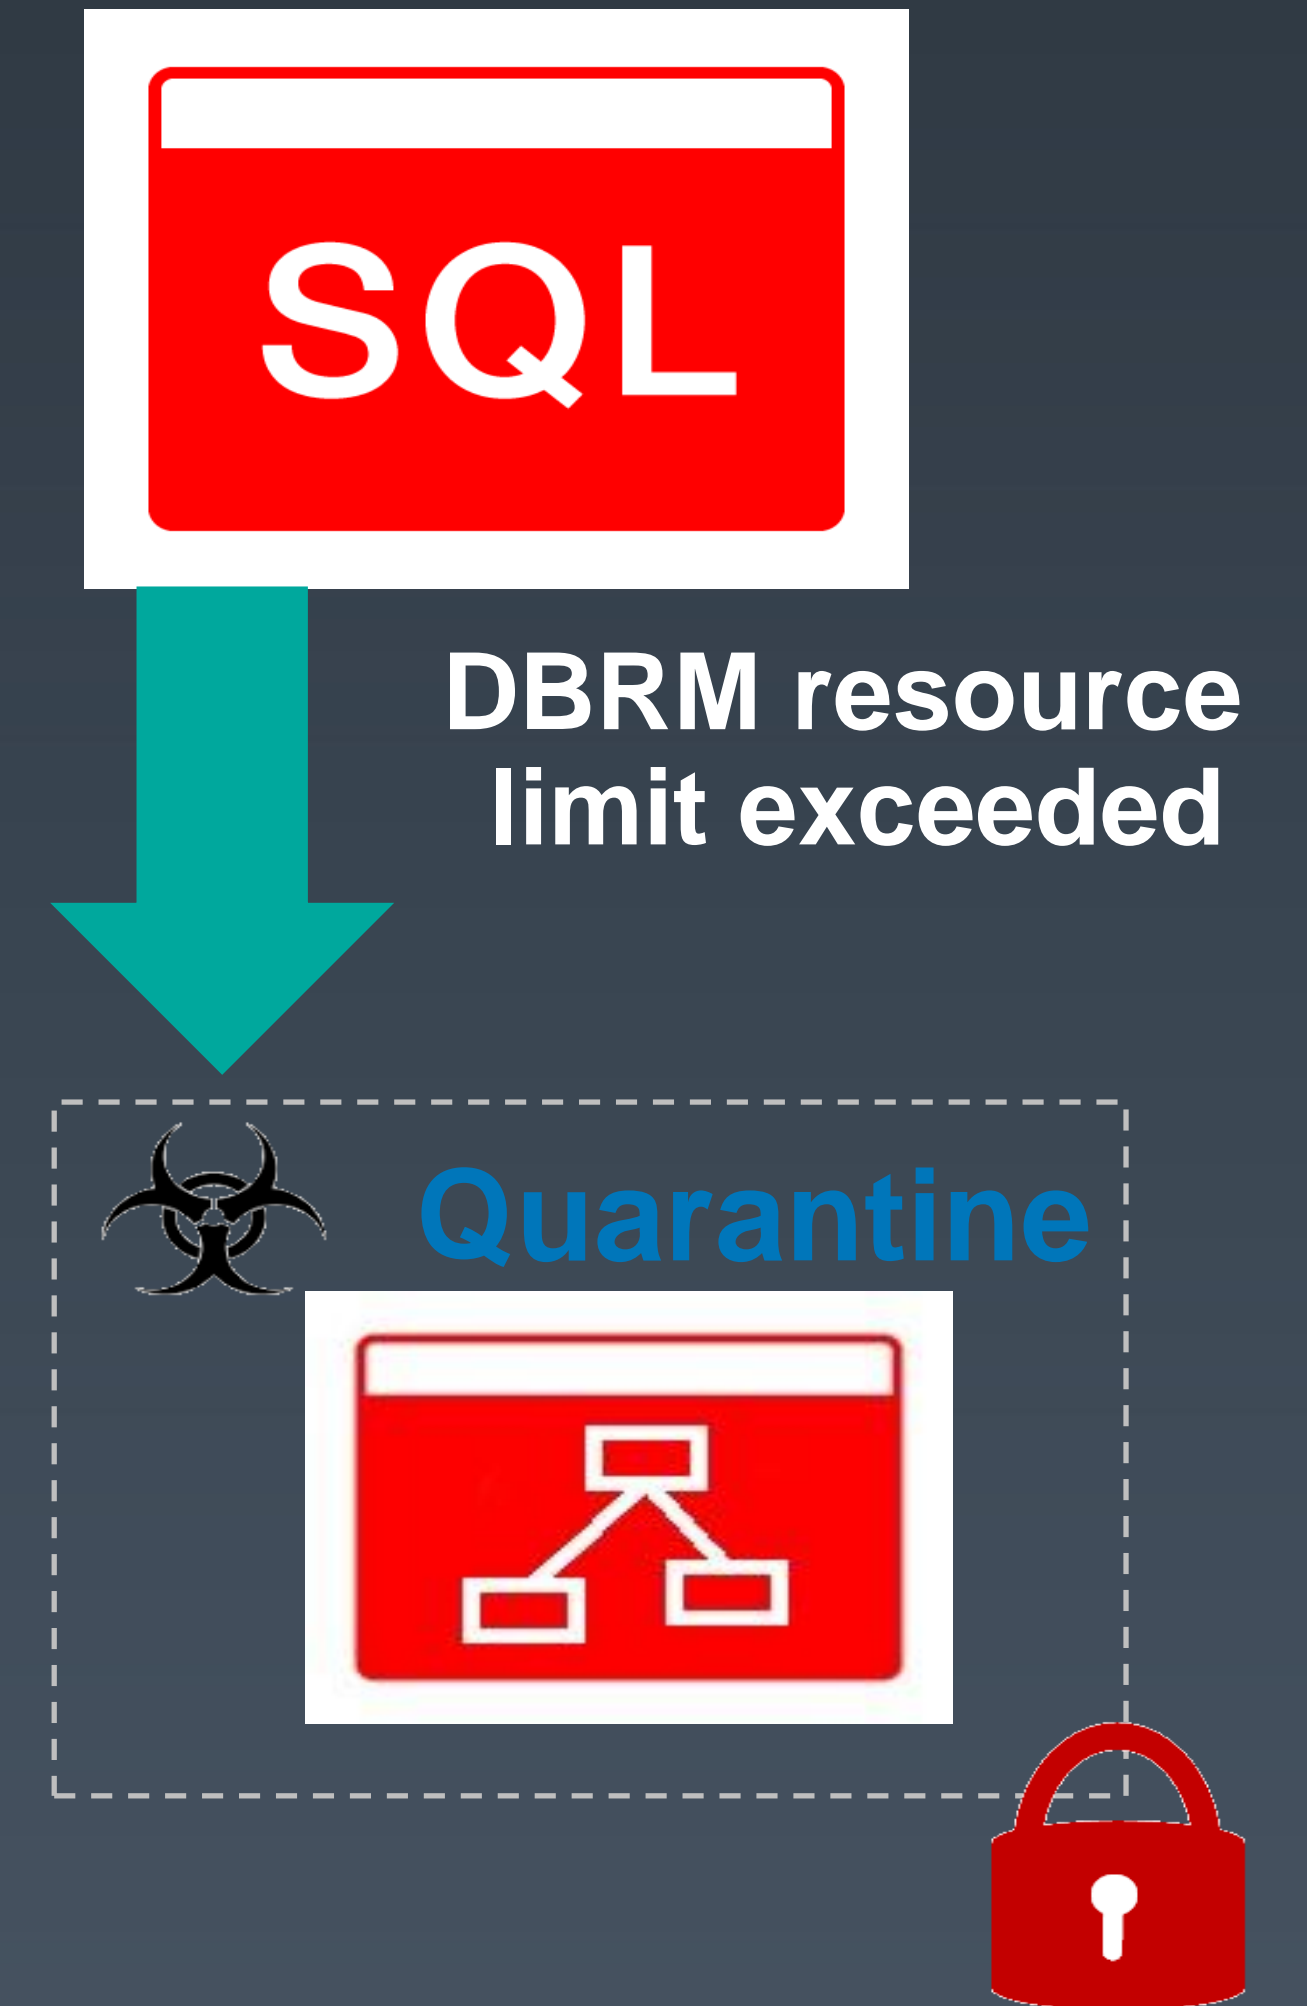
```
SQL> commit;  
Commit complete.
```

Elapsed: **00:00:01.02**

性能管控：失控SQL的自动隔离

NEW IN
19^C

- Oracle Resource Manager 提供SQL监控能力，可以自动终止消耗资源超过一定阈值的SQL。然而，在终止查询之前，大量的资源已经被浪费。
- 在新特性中，执行计划超过DBRM限制将会自动被；
- 隔离的执行计划将被阻止执行；
- SQL 隔离是一个针对失控SQL的自动化解决方案。



性能管控 – 失控计划的隔离

NEW IN
19^C

在以下示例中，SQL 因为执行时间超过限制而被隔离，再次执行时提示“使用了被隔离的执行计划”。

```
SQL> select count(*)
      2  from emp emp1, emp emp2, emp emp3, emp emp4, emp emp5, emp emp6, emp emp7, emp emp8
      3  where rownum <= 10000000;
```

ERROR at line 2:

ORA-00040: **active time limit exceeded - call aborted**

```
SQL> select count(*)
      2  from emp emp1, emp emp2, emp emp3, emp emp4, emp emp5, emp emp6, emp emp7, emp emp8
      3  where rownum <= 10000000;
```

ERROR at line 2:

ORA-56955: **quarantined plan used**

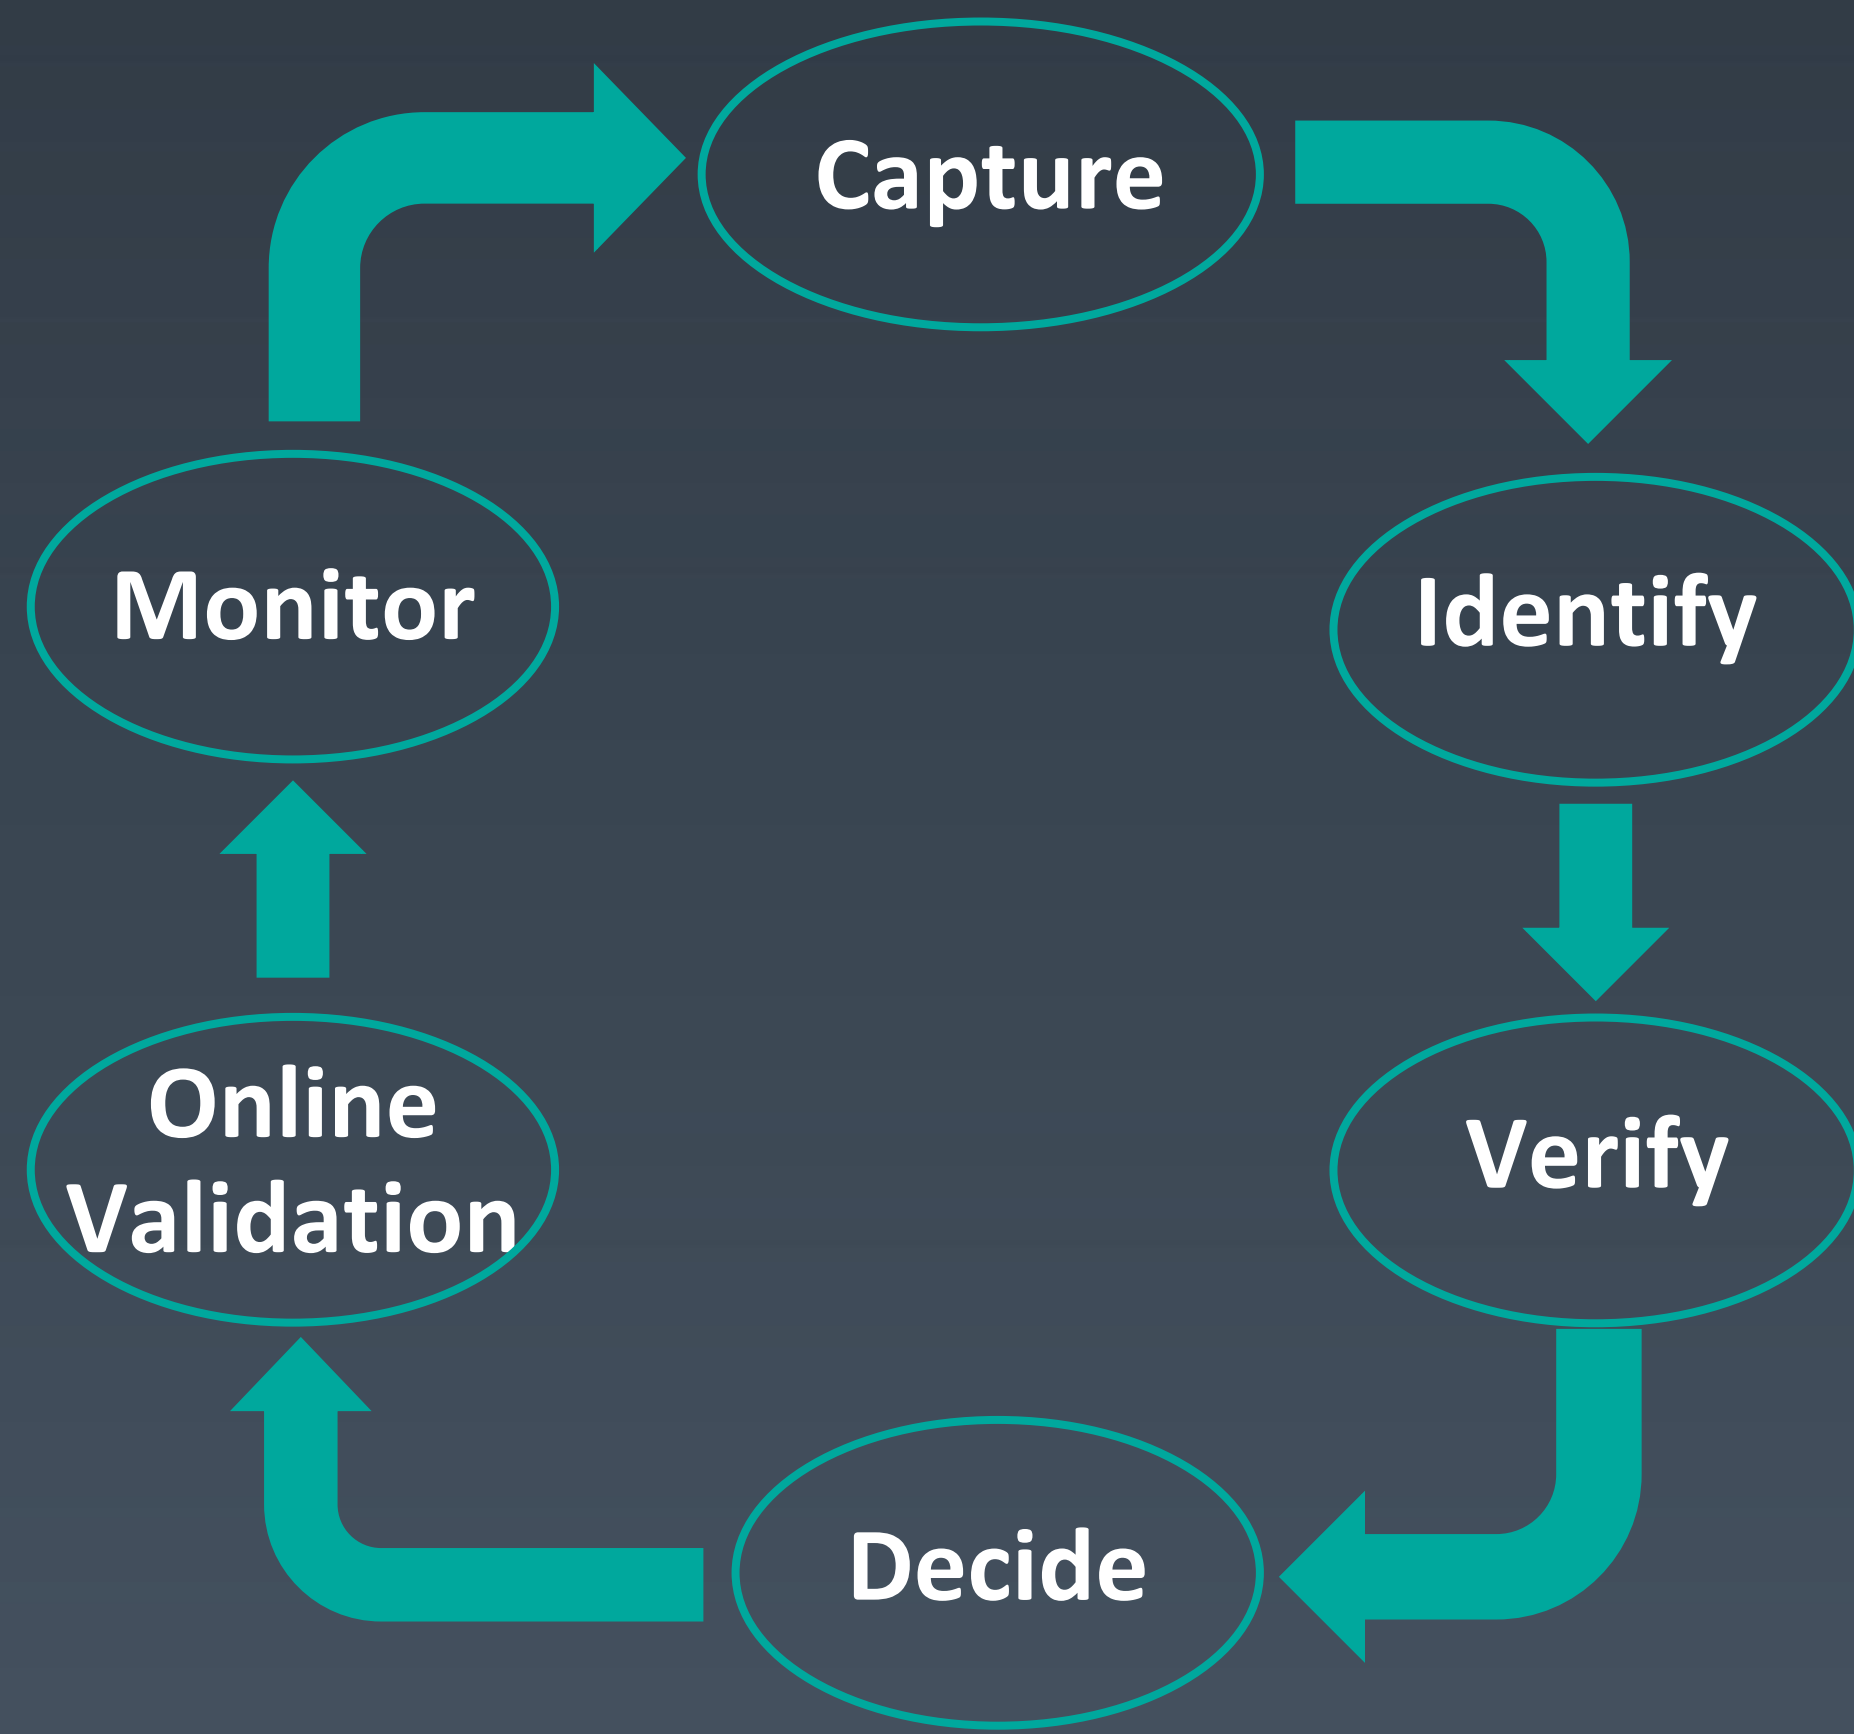
```
SQL> select avoided_executions, sql_quarantine
      2  from v$sql vs
      2  where sql_id = 'd0z9zp1h5n799';
```

| SQL_QUARANTINE | AVOIDED_EXECUTIONS |
|---|--------------------|
| SQL_ QUARANTINE _0scf6as37zcu0cfe7a0e4 | 1 |

自动化索引创建和实施 – Automatic Indexing

NEW IN
19^C

自动索引是借鉴于人工工作的专家系统 It is an **expert system** that implements what a performance engineer skilled in index tuning would do



- 自动索引技术基于和常规手工SQL优化同样的思路实现;
- 系统自动识别候选索引并在启用前验证索引的效率和性能;
- 整个过程完全是自动化实现的;
- 透明度与复杂的自动化同样重要;
 - 所有的调整活动可以通过报告进行核查;

自动化索引技术的实现过程

- 通过 DBA_AUTO_INDEX_CONFIG 修改和启用自动索引特性

```
CDB$ROOT@SYS>EXEC DBMS_AUTO_INDEX.CONFIGURE('AUTO_INDEX_MODE','IMPLEMENT');
CDB$ROOT@SYS>select * from DBA_AUTO_INDEX_CONFIG;
PARAMETER_NAME          PARAMETER_VALUE  LAST_MODIFIED          MODIFIED_BY
-----
AUTO_INDEX_DEFAULT_TABLESPACE
AUTO_INDEX_MODE          IMPLEMENT          17-FEB-19 10.03.59.000000 PM  SYS
AUTO_INDEX_REPORT_RETENTION  31
AUTO_INDEX_RETENTION_FOR_AUTO  373
```

- 通过测试数据执行测试查询，此时表上不存在索引

```
PDB1>create table test as select * from dba_objects;
PDB1>insert into test select * from test;
PDB1>insert into test select * from test;
PDB1>update test set object_id=rownum;
2316704 rows updated.
PDB1>commit;

PDB1>select object_name from test where object_id=1;
PDB1>select object_type from test where object_id=123;
PDB1>select created from test where object_id=345;
```

自动化索引技术的实现过程

- 检查数据库自动任务执行和过程记录

```
select * from DBA_AUTO_INDEX_EXECUTIONS;  
EXECUTION_NAME          EXECUTION_START          EXECUTION_END          ERROR_MESSAGE          ST  
-----  
SYS_AI_2019-02-17/22:51:00      2019-02-17 22:51:00      2019-02-17 22:53:07      CO
```

```
select * from DBA_AUTO_INDEX_STATISTICS where EXECUTION_NAME='SYS_AI_2019-02-17/22:51:00';  
EXECUTION_NAME          STAT_NAME          VALUE  
-----  
SYS_AI_2019-02-17/22:51:00      Index candidates          1  
SYS_AI_2019-02-17/22:51:00      Indexes created (visible)      1  
SYS_AI_2019-02-17/22:51:00      Indexes created (invisible)      0  
SYS_AI_2019-02-17/22:51:00      Indexes dropped          0  
SYS_AI_2019-02-17/22:51:00      Space used in bytes          45088768  
SYS_AI_2019-02-17/22:51:00      Space reclaimed in bytes      0  
SYS_AI_2019-02-17/22:51:00      SQL statements verified      2  
SYS_AI_2019-02-17/22:51:00      SQL statements improved      2  
SYS_AI_2019-02-17/22:51:00      SQL statements managed by SPM      0  
SYS_AI_2019-02-17/22:51:00      SQL plan baselines created      0  
SYS_AI_2019-02-17/22:51:00      Improvement percentage      100
```


自动化索引技术的实现过程

- 以 AI 标识的自动索引已经被创建出来

```
PDB1>select command,statement from DBA_AUTO_INDEX_IND_ACTIONS
      where execution_name='SYS_AI_2019-02-17/22:51:00` order by action_id;
COMMAND          STATEMENT
-----
CREATE INDEX      CREATE INDEX "EN"."SYS_AI_18sc6rdkngxkh" ON "EN"."TEST"("OBJECT_ID")
REBUILD INDEX     ALTER INDEX "EN"."SYS_AI_18sc6rdkngxkh" REBUILD ONLINE
ALTER INDEX VISIBLE ALTER INDEX "EN"."SYS_AI_18sc6rdkngxkh" VISIBLE
```

- 原有查询再次执行将适用智能索引提升性能


```
PDB1>select object_name from test where object_id=1234;
Execution Plan
-----
| Id | Operation | Name | Rows | Bytes | Cost (%CPU) |
-----
| 0 | SELECT STATEMENT | | 1 | 41 | 4 (0) |
| 1 | TABLE ACCESS BY INDEX ROWID BATCHED | TEST | 1 | 41 | 4 (0) |
|* 2 | INDEX RANGE SCAN | SYS_AI_18sc6rdkngxkh | 1 | | 3 (0) |
-----
```

自动化索引技术的实现效果

让完全无索引的应用程序持续运行，
AI系统在后台自动优化创建索引。

运行效果如下，总运行时间 2 小时

- 10 分钟负载高攀；
- 50 分钟负载下降；
- 60 分钟稳定；



总 结

- Oracle 仍然是最好的数据库产品
- 数据库的竞争已经转移到云上
- 数据库的未来是自动化





全球技术领导力峰会

Geekbang | TGO 鲲鹏会
极客邦科技

500+ 高端科技领导者与你一起探讨 技术、管理与商业那些事儿



🕒 2019年6月14-15日 | 📍 上海圣诺亚皇冠假日酒店



扫码了解更多信息

THANKS!

QCon  th